



HUST

Graduation Research 2

Hệ thống Trợ lý Y tế Đa phương thức sử dụng Multi-Agent

Author: Hoàng Quốc Hưng

hung.hq226043@sis.hust.edu.vn

Supervisor: PhD. Dinh Viet Sang

Department: Department of Computer Science

School: School of Information and Communication Technology

Hanoi, Vietnam
January 2026

Mục lục

1	Chapter 1. Introduction	6
1.1	Motivation	6
1.2	Problem Statement	6
1.3	Technical Challenges	6
1.4	Objectives	7
1.5	Contributions	7
1.6	Research Questions	7
1.7	Scope of the Thesis	8
2	Chapter 2. Background and Related Work	8
2.1	Large Language and Vision-Language Models in Healthcare	8
2.2	Medical Question Answering and Multi-hop Reasoning	9
2.3	Retrieval-Augmented Generation for Medical QA	9
2.4	Medical Knowledge Graphs	10
2.5	Multi-agent AI Systems	10
2.6	LoRA and Parameter-Efficient Fine-tuning	11
2.7	Polyp Segmentation and Neoplasm Characterization	11
2.8	Long-Term Memory for AI Agents	12
2.9	Human-in-the-loop for Medical AI	13
3	Chapter 3. Proposed Multi-Agent Medical Assistant Architecture	13
3.1	Overall Architecture	13
3.2	LangGraph-based Agent Orchestration	13
3.3	Agent Cooperation and Cross-module Workflows	15
3.4	Decision Agent	16
3.5	Knowledge and Context Sources	17
3.6	Safety and Human-in-the-loop Validation	17
4	Chapter 4. Text Module: Medical Knowledge Retrieval and Reasoning	17
4.1	Role and Processing Flow of the Text Module	17
4.2	Text Query Types and Routing Behavior	18
4.3	Conversation Agent and Multi-hop QA	19
4.3.1	Illustrative Multi-hop Consultation Flow	19
4.4	Information Sufficiency Assessment	20
4.5	Medical Knowledge Retrieval with RAG and Knowledge Graph	20
4.5.1	Coordinated KG-RAG Evidence Construction	22
4.6	Parallel Reasoning and Fallback Strategy	23
4.7	Safe Medical Response Generation	24
5	Chapter 5. Vision Module: General Medical Image Analysis, Polyp Segmentation, and Polyp VQA	25
5.1	Role and Processing Flow of the Vision Module	25
5.2	Image Routing and Visual Task Classification	26
5.3	General Medical Image Agent	26
5.4	Polyp Segmentation Agent	27
5.5	DeepLabV3+ for Polyp Segmentation	27
5.6	Post-processing and Overlay Mask	29
5.7	Polyp VQA Agent Design with LoRA Fine-tuned VLMs	32
5.8	Segmentation-to-VQA Agent Cooperation	32

5.8.1	Example Interaction Pattern	34
5.9	Original-image and Segmentation-grounded VQA Settings	34
6	Chapter 6. Memory Layer: Long-Term Context for Medical Agents	35
6.1	Role of the Memory Layer	35
6.2	Integration Architecture	35
6.3	Technical Pipeline	35
6.4	Impact on Agent Reasoning	36
6.5	Safety Constraints and Limitations	36
7	Chapter 7. Data and Experimental Setup	36
7.1	VimedAQA Dataset	36
7.2	VietMedKG	37
7.3	BKAI-IGH NeoPolyp-Small Dataset	38
7.4	Kvasir-VQA-X1 Dataset	39
7.5	Polyp VQA Fine-tuning Setup	40
7.6	Memory Evaluation Scenarios	41
7.7	Evaluation Metrics	41
7.8	Experimental Protocol	41
8	Chapter 8. Results and Analysis	42
8.1	Text Module Results and Analysis	42
8.2	Polyp Segmentation Results and Analysis	43
8.3	Polyp VQA Fine-tuning Results and Analysis on Kvasir-VQA-X1	44
9	Chapter 9. Overall Discussion and Limitations	46
9.1	Strengths of the Proposed System	46
9.2	Text Module Limitations	46
9.3	Vision Module Limitations	47
9.4	Kvasir-VQA-X1 Limitations and Domain Transfer	47
9.5	Memory Layer Limitations	47
9.6	Evaluation and Safety Limitations	48
10	Chapter 10. Conclusion and Future Work	48
10.1	Summary of Achievements	48
10.2	Technical Contributions	48
10.3	Future Work	48

List of Tables

1	Benefits of the multi-agent design illustrated by the detailed workflow	15
2	Text query types and corresponding routing behavior	18
3	Coordinated evidence construction in the Text Module	23
4	Complementary roles of the Polyp Segmentation Agent and Polyp VQA Agent	33
5	Local VimedAQA record distribution used for RAG indexing	37
6	Kvasir-VQA-X1 training data before and after segmentation-mask filtering	39
7	LoRA fine-tuning hyperparameters for Polyp VQA	40
8	Text Module benchmark results	42
9	Dice Score comparison for polyp segmentation models	43
10	Polyp VQA evaluation on Kvasir-VQA-X1	45

11	LLM Judge comparison between LoRA-only and visual-grounded MedGemma on the mask-matched subset	45
----	--	----

List of Figures

1	Overall architecture of the Multi-Agent Medical Assistant	13
2	Detailed agent workflow and module coordination	14
3	Text Module architecture for medical knowledge retrieval	21
4	Parallel KG-RAG processing in the Text Module	21
5	Medical chain-of-thought and information sufficiency flow	24
6	Vision Module workflow for medical image analysis	26
7	DeepLabV3+ encoder-decoder architecture with Atrous Spatial Pyramid Pooling	28
8	Visualization of raw and post-processed polyp segmentation masks	31
9	Medical QA dataset overview used for retrieval-oriented experiments	37
10	Overview of the VietMedKG data source used by the Knowledge Graph Agent .	38
11	Example images and annotations from the BKAI-IGH NeoPolyp-Small dataset .	38
12	Example structure of the Kvasir-VQA-X1 dataset for gastrointestinal endoscopy VQA	39

Preface

The increasing demand for reliable, accessible, and personalized medical information motivates the development of intelligent medical assistant systems. Users often describe symptoms in incomplete or informal language, search for medical information from sources of uneven quality, and may need support across multiple interactions rather than a single isolated conversation.

Recent advances in large language models, retrieval-augmented generation, knowledge graphs, multimodal models, and agent orchestration provide an opportunity to build medical AI systems that are more grounded, modular, and context-aware. However, medical AI also introduces safety challenges, including hallucination, incomplete clinical context, uncertain visual interpretation, and the need for human feedback.

This report presents a Multi-Agent Medical Assistant that integrates text-based medical reasoning, medical image analysis, long-term memory, and human-in-the-loop validation. The work focuses on AI techniques and system-level reasoning rather than user interface or backend implementation details.

Summary of Content and Contributions

This report studies a multi-agent medical assistant designed around four main technical pillars:

- **Text-based medical reasoning:** The system combines Multi-hop Question Answering, Retrieval-Augmented Generation, Knowledge Graph reasoning, and Web Search fallback to produce more grounded and cautious medical responses.
- **Vision-based medical reasoning:** The Vision Module supports general medical image interpretation, colonoscopy polyp segmentation, and Polyp Visual Question Answering. The Polyp VQA component uses a LoRA fine-tuned Vision-Language Model and compares original-image reasoning with segmentation-grounded reasoning.
- **Long-term memory for agent continuity:** Mem0 is integrated as a memory layer that stores clinically useful patient-reported facts across sessions, helping agents maintain continuity while avoiding unsupported diagnostic claims.
- **Human-in-the-loop validation:** The system includes a validation mechanism for collecting user feedback and supporting safer review of generated responses.

The main contribution of this work is an AI-centered architecture that combines specialized agents, structured and unstructured knowledge retrieval, visual grounding, long-term memory, and validation-oriented safety mechanisms for medical assistance.

1 Chapter 1. Introduction

1.1 Motivation

Healthcare information seeking is increasingly mediated by digital systems, yet most users still face a fundamental gap between their symptoms and reliable medical interpretation. A patient may describe a symptom in informal language, omit clinically important details, misunderstand medication risks, or combine several concerns across different interaction sessions. At the same time, medical information available on the internet is abundant but uneven in quality, difficult to verify, and often not adapted to the user’s clinical context.

Recent progress in large language models (LLMs), retrieval-augmented generation (RAG), knowledge graphs, and vision-language models (VLMs) makes it possible to design medical assistants that can reason over both text and images. However, a direct application of a single generative model is not sufficient for a safety-critical domain. Medical consultation requires source grounding, uncertainty awareness, follow-up questioning, multimodal interpretation, and mechanisms for human validation.

This thesis develops a **Multi-Agent Medical Assistant** that combines text-based medical reasoning, medical image analysis, long-term memory, and human-in-the-loop validation. The system is designed as a coordinated set of specialized agents rather than a single monolithic model. This design allows each subtask to be handled by a component with a suitable knowledge source or model capability.

1.2 Problem Statement

The problem addressed in this thesis is the design and evaluation of a multi-agent medical assistant that can:

- answer medical questions using grounded knowledge sources;
- collect missing clinical information through multi-hop questioning;
- combine RAG, Knowledge Graph reasoning, and web search for text-based medical reasoning;
- analyze medical images, especially colonoscopy polyp images;
- perform polyp segmentation and use segmentation masks as visual grounding for VQA;
- maintain clinically useful long-term patient context through a memory layer;
- support human-in-the-loop validation for safer output review.

The focus is not on building a web application interface, but on the AI architecture, reasoning workflow, multimodal coordination, memory integration, and experimental analysis.

1.3 Technical Challenges

Several challenges motivate the proposed architecture.

Incomplete clinical context. Users often provide insufficient symptom descriptions. A reliable medical assistant should recognize when information is missing and ask follow-up questions before generating a recommendation.

Hallucination and source reliability. LLMs can generate fluent but unsupported medical content. Grounding responses in RAG documents, Knowledge Graph facts, and trusted external sources reduces this risk.

Structured and unstructured medical knowledge. Medical reasoning often requires both detailed textual explanation and structured relations between symptoms, diseases, drugs, and risk factors. RAG and Knowledge Graphs provide complementary capabilities.

Medical image understanding. Medical image analysis requires models that can identify

visual patterns and communicate findings carefully. In colonoscopy images, polyp localization and neoplasm characterization are important subtasks.

Visual grounding for VQA. A VLM answering questions from an image may attend to irrelevant regions. Segmentation masks can guide the VQA model toward clinically relevant regions, but incorrect masks may also introduce noise.

Long-term patient context. Medical interaction is rarely isolated. Information such as allergies, medication use, chronic diseases, or repeated symptoms may remain relevant across sessions. A memory layer must preserve useful facts without treating them as official diagnoses.

Safety and validation. Medical AI systems require feedback mechanisms and explicit boundaries. Human-in-the-loop validation helps collect feedback and identify responses that may require review.

1.4 Objectives

The main objective is to design and evaluate an AI-centered multi-agent medical assistant with four core capabilities:

- **Text Module:** medical reasoning with Conversation Agent, Multi-hop QA, RAG, Knowledge Graph, Web Search, and safe response generation;
- **Vision Module:** general medical image interpretation, polyp segmentation, and Polyp VQA;
- **Memory Layer:** long-term patient context using Mem0-style memory integration;
- **Human Validation:** user validation and feedback logging as a human-in-the-loop safety mechanism.

1.5 Contributions

The main contributions of this thesis are:

- a multi-agent architecture for medical reasoning and multimodal consultation;
- a Text Module that combines Multi-hop QA, RAG, Knowledge Graph reasoning, and fallback search;
- a Vision Module that integrates General Medical Image Agent, Polyp Segmentation Agent, and Polyp VQA Agent;
- a polyp segmentation pipeline with post-processing based on connected components and region-wise color dominance refinement;
- a Polyp VQA approach using LoRA fine-tuned MedGemma-4B-it and Qwen3.5-4B settings, evaluated with original image input and visual grounding with segmentation mask;
- a long-term memory layer using Mem0 for patient-context continuity;
- an evaluation design covering Dice Score, BLEU, ROUGE-L, LLM Judge, One-Hop Accuracy, Multi-Hop Awareness Score, and memory-oriented system behavior.

1.6 Research Questions

This thesis is guided by five research questions. First, how can a medical assistant combine conversational reasoning, RAG, Knowledge Graph reasoning, and Web Search to reduce unsupported medical responses? Second, how can information sufficiency assessment improve safety in symptom-based consultation? Third, can segmentation masks improve Polyp VQA by providing visual grounding for a fine-tuned VLM? Fourth, how can long-term memory support personalization without turning patient-reported facts into verified diagnoses? Finally, what trade-offs appear between modularity, response quality, latency, and safety in a multi-agent medical assistant?

These questions connect the architecture with the evaluation methodology. The report therefore evaluates the system as a reasoning architecture and not as a clinical product ready for deployment.

1.7 Scope of the Thesis

The thesis focuses on AI methods and system-level reasoning. Those components are necessary for system integration, but the main contribution lies in agent orchestration, knowledge retrieval, multimodal reasoning, long-term memory, and validation-oriented safety.

The medical scope is also limited. The system is intended to provide information support and preliminary reasoning. It does not replace physicians, does not provide definitive diagnosis, and does not make treatment decisions. For high-risk symptoms or uncertain findings, the system should recommend professional medical consultation.

2 Chapter 2. Background and Related Work

2.1 Large Language and Vision-Language Models in Healthcare

Large language models have become a central direction in medical AI because they can represent clinical knowledge, answer medical questions, and generate patient-facing explanations [10]. Med-PaLM and Med-PaLM 2 show that instruction tuning, medical-domain adaptation, and prompting strategies can substantially improve medical question answering [11]. Open medical LLMs such as PMC-LLaMA and BioMistral follow another direction by adapting general-purpose models to biomedical literature, textbooks, and medical instruction data [12, 13]. These studies show clear progress, but they mostly evaluate the model as an independent QA system rather than as part of a complete consultation workflow.

Medical vision-language models extend this trend from text-only reasoning to image-text reasoning. LLaVA-Med and Med-Flamingo adapt multimodal foundation models for biomedical image conversation, VQA, and rationale generation [14, 15]. More recent models such as MedGemma provide medical vision-language backbones, while general model families such as Qwen offer strong multilingual and instruction-following capabilities that can be adapted to medical tasks [38, 39]. In gastrointestinal endoscopy, datasets such as Kvasir-VQA and Kvasir-VQA-x1 further indicate a shift from image classification toward question answering over clinical visual evidence [35, 36].

However, current LLMs and VLMs still have important limitations for medical assistants. They can generate fluent but unsupported answers, especially when the user question lacks clinical context or when the model relies only on parametric memory [16]. Many models also operate without explicit retrieval, Knowledge Graph evidence, segmentation grounding, or human validation. For VLMs, correct-looking answers may come from dataset priors or textual shortcuts rather than from the relevant image region, which is risky in endoscopy because lesion interpretation depends on local boundary, texture, morphology, and color patterns.

This thesis therefore treats LLMs and VLMs as components inside a larger safety-aware architecture. The proposed system combines model generation with routing, information sufficiency assessment, RAG, Knowledge Graph reasoning, memory, segmentation-based visual grounding, and human validation. In the Polyp VQA setting, MedGemma-4B-it and Qwen3.5-4B are adapted with LoRA and evaluated under both original-image and segmentation-grounded inputs. This design directly addresses the limitation of standalone models by testing whether external evidence and visual grounding improve reasoning reliability rather than only answer fluency.

2.2 Medical Question Answering and Multi-hop Reasoning

Medical question answering has traditionally been evaluated through benchmark datasets such as PubMedQA, MedQA, and MedMCQA [17, 18, 19]. These datasets are valuable because they test biomedical reading comprehension, medical exam knowledge, and multiple-choice clinical reasoning. Vietnamese medical QA datasets such as VimedAQA further show the importance of language-specific resources for local healthcare applications [6]. However, these benchmarks usually provide a well-formed question and a fixed answer space. Real users often provide incomplete, informal, or ambiguous descriptions, so benchmark performance alone is not sufficient to evaluate a consultation-oriented assistant.

The key difference is that medical consultation is frequently underspecified at the first turn. A user asking about chest pain, abdominal pain, fever, dizziness, or medication use may omit onset time, severity, duration, associated symptoms, age, pregnancy status, allergies, chronic diseases, or current medication. A direct answer may therefore be clinically incomplete even if it is linguistically fluent. This creates a gap between single-turn medical QA and interactive medical reasoning.

Multi-hop reasoning addresses this gap by decomposing a complex question into intermediate reasoning steps and evidence requirements. General multi-hop QA research shows that answering some questions requires connecting multiple supporting facts rather than retrieving one isolated passage [20]. In the medical domain, the same principle appears in two forms. The first form is evidence multi-hop reasoning, where the system combines information from multiple sources such as retrieved documents, Knowledge Graph relations, and previous conversation context. The second form is dialogue multi-hop reasoning, where the system asks follow-up questions before generating a final response.

This thesis focuses on the second form as a safety mechanism and uses the first form for evidence synthesis. The Conversation Agent first evaluates whether the available user information is sufficient. If important variables are missing, the system asks targeted follow-up questions instead of producing a premature conclusion. When sufficient context is available, the Text Module can route the query to RAG, Knowledge Graph reasoning, web search, or response generation. This workflow makes multi-hop QA more than a prompting technique: it becomes a control mechanism for deciding when to answer, when to retrieve, when to ask, and when to recommend professional care.

Multi-hop reasoning is also necessary for separating general medical education from patient-specific advice. A question such as “What is gastritis?” can often be answered using retrieved explanatory material. In contrast, a question such as “Do I have gastritis?” requires symptom duration, pain location, triggers, warning signs, medication history, and possibly physician evaluation. The proposed Text Module uses this distinction to reduce unsafe one-shot answers. In this sense, multi-hop QA provides the bridge between medical LLM capability and the thesis objective of building a safer multi-agent medical assistant.

2.3 Retrieval-Augmented Generation for Medical QA

Retrieval-Augmented Generation combines neural generation with external document retrieval [7]. In a medical assistant, RAG can reduce hallucination by grounding the model response in retrieved medical documents. A typical RAG pipeline includes query expansion, embedding-based retrieval, document reranking, and response generation with source context.

The proposed system uses RAG for Vietnamese medical question answering. Documents are embedded and stored in a vector database, while the response generator uses retrieved context to produce structured answers. This approach is useful when the answer requires explanatory text, treatment descriptions, disease overviews, or patient-friendly medical information.

RAG is not sufficient by itself. Retrieval can fail when a query is underspecified, when relevant documents are absent from the corpus, or when semantically similar but clinically irrelevant passages are retrieved. For this reason, the proposed Text Module combines RAG with query expansion, reranking, context filtering, and Knowledge Graph reasoning.

Medical RAG also requires attention to language and domain mismatch. In Vietnamese medical QA, users may mix colloquial expressions with formal medical terms. A retrieval system must bridge this gap without retrieving unrelated passages. Query expansion and filtering are therefore not auxiliary features; they are necessary parts of making retrieval usable in real consultation-style interactions.

2.4 Medical Knowledge Graphs

Knowledge Graphs represent medical knowledge as entities and relations. Entities may include diseases, symptoms, drugs, body parts, risk factors, tests, or treatments. Relations encode structured links such as disease-has-symptom, drug-treats-disease, or disease-associated-with-risk-factor.

Compared with RAG, a Knowledge Graph is more suitable for structured reasoning. It can support explicit traversal of disease-symptom relationships and provide more interpretable evidence for certain diagnostic suggestions. In this thesis, the Knowledge Graph Agent queries VietMedKG and supports reasoning for symptom-to-disease and related medical questions.

Knowledge Graphs also have limitations. They depend on relation coverage and graph construction quality. A graph may encode important disease-symptom relations but lack detailed explanation or updated guidance. Therefore, this thesis treats KG and RAG as complementary rather than competing approaches.

The complementarity is especially relevant for medical reasoning. A Knowledge Graph can provide a compact relation such as a symptom-disease link, but a patient-facing answer often requires explanation, uncertainty, and practical next steps. RAG can provide that explanatory material. Conversely, RAG may retrieve long text but fail to make an explicit relation clear. KG can provide structure to support the final synthesis.

2.5 Multi-agent AI Systems

Multi-agent systems divide complex tasks among specialized agents. In medical AI, this is useful because different queries require different capabilities. A symptom consultation may require multi-hop questioning, a disease explanation may require RAG, a structured reasoning query may require Knowledge Graph traversal, and an image question may require a vision-language model.

The proposed system uses a Decision Agent to route user requests to the appropriate agents. The workflow is orchestrated with LangGraph [1], which allows the system to maintain state, coordinate agent execution, and route to fallback strategies when internal knowledge is insufficient.

The multi-agent design also improves observability. The system can identify which agent handled a query and which source of information was used. This is useful for debugging, evaluation, and user trust in safety-sensitive medical applications.

However, multi-agent systems introduce coordination challenges. Routing errors, inconsistent outputs between agents, and increased latency can reduce system quality. A medical multi-agent assistant must therefore include clear routing rules, fallback behavior, and evaluation of intermediate decisions. This thesis treats agent orchestration as a core technical contribution rather than as an implementation detail.

2.6 LoRA and Parameter-Efficient Fine-tuning

Fine-tuning large VLMs is computationally expensive because full-parameter updates require significant memory and training resources. LoRA provides a parameter-efficient alternative by learning low-rank adaptation matrices while keeping most original model parameters fixed [37]. This reduces the training cost and makes domain adaptation more practical.

In this thesis, LoRA is used to adapt VLMs for Polyp VQA. The goal is to improve the model’s ability to answer endoscopy-related visual questions while keeping fine-tuning feasible under limited resources. This is appropriate for the thesis setting because endoscopic VQA data is specialized and smaller than the generic multimodal corpora used to train foundation models.

The fine-tuning experiment is designed around two compact base models: MedGemma-4B-it and Qwen3.5-4B. MedGemma is a medical vision-language model based on the Gemma family and is designed to provide stronger medical image and text understanding than a general-purpose model of similar scale [38]. Qwen models provide a strong multilingual and instruction-following foundation, with Qwen3 extending the family toward hybrid reasoning and broader multilingual coverage [39]. Comparing these two backbones is useful because Polyp VQA requires both visual understanding and language generation: MedGemma contributes domain-specific medical pretraining, while Qwen provides a strong general instruction-following baseline.

LoRA also makes it practical to compare different input designs. In this work, the same VQA task is evaluated with original-image input and segmentation-grounded input, allowing the analysis to separate the effect of model adaptation from the effect of visual grounding. The comparison should therefore report results for each base model and each input setting: MedGemma-4B-it with original image, MedGemma-4B-it with segmentation grounding, Qwen3.5-4B with original image, and Qwen3.5-4B with segmentation grounding.

Parameter-efficient fine-tuning is particularly useful in medical VQA because the available domain-specific data is often smaller than general vision-language datasets. Instead of training a large model from scratch, LoRA adapts the model toward the target domain while preserving most of the base model’s general multimodal ability.

2.7 Polyp Segmentation and Neoplasm Characterization

Colonoscopy polyp segmentation is a key medical image analysis task because accurate localization of polyps can support colorectal cancer screening and clinical interpretation. The task can be binary, where the model separates polyp from background, or multi-class, where the model also distinguishes neoplastic and non-neoplastic regions. Compared with general semantic segmentation, polyp segmentation is difficult because polyps vary greatly in size, color, texture, shape, and boundary clarity. The boundary between polyp tissue and surrounding mucosa can be ambiguous, and artifacts such as blur, specular highlights, low contrast, and imperfect bowel preparation can further reduce visual quality [30].

Public datasets have strongly shaped research in this area. Kvasir-SEG provides manually annotated gastrointestinal polyp images with segmentation masks and is widely used for benchmarking binary polyp segmentation [31]. For neoplasm characterization, the BKAI-IGH NeoPolyp dataset extends the problem from localization to multi-class segmentation, where the model must distinguish neoplastic and non-neoplastic regions. This makes the task closer to clinical interpretation than a pure foreground-background mask, but it also increases class imbalance and label ambiguity.

Early deep learning approaches for biomedical segmentation were strongly influenced by U-Net, whose encoder-decoder structure combines contextual feature extraction with precise localization through skip connections [27]. U-Net++ further redesigned skip pathways with

nested dense connections to reduce the semantic gap between encoder and decoder features [28]. These architectures remain important baselines because medical datasets are often small and require models that can learn from limited annotated data.

DeepLabV3+ introduced another influential direction by combining atrous convolution, Atrous Spatial Pyramid Pooling, and a decoder for boundary refinement [29]. The multi-scale receptive fields are useful for polyp segmentation because lesions can occupy very different image areas. This thesis follows the DeepLabV3+ family for the deployed segmentation model and produces a three-class output. In the BKAI-IGH NeoPolyp-Small setting, DeepLabV3+ with RegNetx320 achieves the strongest Dice Score after post-processing among the evaluated models.

Attention-based and transformer-based methods address different limitations. PraNet uses a parallel partial decoder and reverse attention to refine boundary cues and recover missed regions [30]. Polyp-PVT replaces a CNN backbone with a pyramid vision transformer and introduces modules for cascaded fusion, camouflage identification, and similarity aggregation, improving robustness in challenging cases such as small objects and appearance variation [32]. More recent medical segmentation work also emphasizes efficient decoders; EMCAD proposes a multi-scale convolutional attention decoder that reduces parameters and FLOPs while maintaining strong performance across medical segmentation tasks [33].

Foundation segmentation models introduce another research direction. Polyp-SAM adapts the Segment Anything Model to the polyp domain and shows that foundation models can be transferred to medical segmentation, but domain adaptation remains necessary because generic segmentation priors do not automatically solve endoscopic boundary ambiguity [34]. This supports the broader design principle of the thesis: general foundation models are useful starting points, but task-specific adaptation and evaluation are still required for medical workflows.

The segmentation output is not only an end result. It is also used as a bridge between computer vision and multimodal reasoning. A clean mask can help the Polyp VQA model focus on clinically relevant regions and provide a more grounded answer. This is where the proposed work differs from studies that evaluate segmentation alone. The segmentation agent produces spatial evidence, the VQA agent converts image-question inputs into language, and LoRA fine-tuning adapts MedGemma-4B-it and Qwen3.5-4B backbones to the endoscopic VQA task. The central hypothesis is that segmentation masks can improve VQA when the mask correctly localizes clinically relevant tissue, but they can also harm the answer if the mask is wrong or visually misleading. Therefore, segmentation quality, visual grounding, and VQA quality must be evaluated together rather than as isolated components.

2.8 Long-Term Memory for AI Agents

Most dialogue agents maintain short-term state within a session, but medical interaction often requires continuity across sessions. Long-term memory allows an agent to remember clinically useful facts such as allergies, chronic diseases, medication use, repeated symptoms, and patient preferences.

In the proposed system, Mem0 is used as a long-term memory layer. It does not replace LangGraph, RAG, the Knowledge Graph, or patient data storage. Instead, it provides additional context that helps agents personalize responses and avoid repeatedly asking for information that was already provided.

Medical memory must be handled conservatively. A stored fact may represent a user-reported allergy or symptom history, but it should not be treated as a verified diagnosis. Memory is therefore used as supporting context rather than authoritative clinical truth.

2.9 Human-in-the-loop for Medical AI

Human-in-the-loop validation is important in medical AI because user feedback can reveal unsafe, incomplete, or unclear responses. In this system, validation allows users to confirm or comment on system responses. This feedback is logged and can support quality review, memory selection, and future improvement.

3 Chapter 3. Proposed Multi-Agent Medical Assistant Architecture

3.1 Overall Architecture

The proposed system is organized as a multi-agent medical assistant with four main layers:

- **Decision and orchestration layer:** selects the appropriate processing path;
- **Text Module:** handles medical QA, multi-hop consultation, RAG, Knowledge Graph, and fallback search;
- **Vision Module:** handles general medical image interpretation, polyp segmentation, and Polyp VQA;
- **Memory and validation layer:** provides long-term patient context and human-in-the-loop feedback.

This architecture is designed to reduce reliance on a single model. Instead of asking one LLM to solve every task, the system delegates work to agents that are specialized for retrieval, graph reasoning, visual interpretation, segmentation, or VQA.




Figure 1: Overall architecture of the Multi-Agent Medical Assistant

Figure 1 shows that the system is not organized as a linear chatbot pipeline. The Decision Agent first interprets the modality and intent of the user input, then routes the request to the Text Module or Vision Module. Shared memory is placed between these modules rather than attached to only one agent. This design allows text-based reasoning and image-based reasoning to exchange context while still keeping their internal models and evidence sources separate.

The high-level architecture in Figure 1 establishes the main separation of responsibilities. The more detailed workflow in Figure 2 explains how these responsibilities are coordinated during execution and is used below to analyze the system benefits.

3.2 LangGraph-based Agent Orchestration

LangGraph is used as the workflow orchestration framework. It maintains state across the processing flow and connects decision nodes with task-specific agents. The graph-based struc-

ture is useful because medical consultation is not a simple single-step process. A query may require routing, follow-up questioning, retrieval, fallback search, image analysis, validation, or memory update.

The short-term state includes the current query, image path when available, selected agent, conversation history, patient identifier, and intermediate outputs such as segmentation paths or retrieved contexts. This state supports coherent routing and response generation within the current session.

The graph structure also makes the workflow easier to control than a purely prompt-based pipeline. Each node has a clear role: routing, retrieval, image analysis, response generation, validation, or memory update. This is useful in a medical context because the system can inspect intermediate outputs before producing the final answer. For example, if retrieval fails or the image route is uncertain, the workflow can move to a fallback path instead of forcing a response from incomplete evidence.

Another advantage is that LangGraph separates short-term state from long-term memory. Short-term state is useful for the current interaction, while Mem0 stores facts that may remain useful across sessions. This separation prevents temporary reasoning artifacts, such as a retrieved passage or a provisional model output, from automatically becoming persistent patient memory.




Figure 2: Detailed agent workflow and module coordination

Figure 2 makes the advantages of the multi-agent design easier to observe because it shows not only the system components, but also the routes between them. The benefits are summarized in Table 1. This presentation separates the benefit, the corresponding workflow evidence, and its implication for a medical assistant.

Benefit	How Figure 2 supports the benefit
Functional decomposition	User input is routed from the Agent Core to specialized text and vision branches. This means heterogeneous medical tasks can be handled by agents with suitable tools instead of forcing one model to perform retrieval, reasoning, segmentation, and VQA at the same time.
Evidence specialization	KG, RAG, Web Search, shared memory, and segmentation masks appear as separate evidence channels. The system can therefore combine structured relations, explanatory documents, updated web evidence, patient context, and visual grounding without collapsing them into one opaque prompt.
Graceful fallback	Information Assessment routes successful cases to MedGemma CoT and weak-evidence cases to Web Search or conversational handling. This helps avoid premature answers when internal evidence is missing, uncertain, or clinically underspecified.
Cross-modal continuity	Vision outputs and text reasoning interact through shared memory and workflow state. A user can upload an image, receive an initial interpretation, and then continue asking follow-up questions that reuse the image-derived context.
Observability and evaluation	Intermediate modules such as routing, query expansion, KG/RAG retrieval, segmentation, and response generation are explicit. Errors can therefore be traced to a specific stage, and the system can be evaluated beyond the final generated answer.

Table 1: Benefits of the multi-agent design illustrated by the detailed workflow

This table also clarifies why the proposed system is not merely a collection of independent models. The value of the architecture comes from how intermediate evidence is produced, assessed, reused, and routed. For example, the Polyp Segmentation Agent produces a spatial mask that can later guide Polyp VQA, while the KG and RAG agents provide complementary text evidence for the final response. The detailed workflow therefore provides a concrete basis for discussing system benefits rather than presenting them as abstract claims.

3.3 Agent Cooperation and Cross-module Workflows

A key characteristic of the proposed system is that agents cooperate through shared workflow state, memory, and intermediate evidence objects. In this design, the output of one agent can become a structured input for another agent. This is important because many medical-assistant tasks require more than one capability: a system may need to detect the modality of the request, retrieve medical evidence, localize a visual region, ask follow-up questions, and synthesize a cautious response in the same interaction.

The first cooperation pattern appears inside the Text Module. The Conversation Agent, Query Expansion component, KG Agent, RAG Agent, Web Search Agent, and response generator form a coordinated reasoning pipeline. The Conversation Agent determines whether the user has provided enough clinical context. If the request is medical and sufficiently specified, Query Expansion reformulates the query into medically useful terms. The KG Agent then provides structured relations such as symptom–disease links, while the RAG Agent retrieves explanatory documents from the vector database. Information Assessment evaluates whether these sources are sufficient; if not, the Web Search Agent is activated as a fallback. The final answer is therefore not produced from a single source, but from a controlled cooperation among agents with complementary evidence roles.

The second cooperation pattern appears between the Vision Module and the Polyp VQA Agent. The Polyp Segmentation Agent produces a mask that localizes clinically relevant regions in a colonoscopy image. This mask is not only a final segmentation output; it also becomes a visual grounding signal for the Polyp VQA Agent. The VQA Agent can then answer natural-language questions by using both the original image and the segmentation overlay. This cooperation combines the precision of pixel-level localization with the reasoning ability of a fine-tuned vision-language model. The experimental comparison between original-image VQA and segmentation-grounded VQA evaluates whether this cooperation improves answer quality, rather than assuming that agent collaboration is beneficial without evidence.

The third cooperation pattern is cross-modal interaction between the Text Module and Vision Module. After a user uploads an image and receives an initial visual interpretation, the conversation does not need to end. The user can continue asking follow-up questions such as asking what a highlighted region means, whether the finding is concerning, what additional symptoms should be monitored, or what medical consultation may be appropriate. The system can keep the image-derived context in short-term state and route later text questions back to the relevant visual evidence or to the Text Module for medical explanation. This makes the assistant closer to an interactive consultation workflow than a one-shot image classifier.

Shared memory further supports this cooperation by preserving clinically useful context across turns and sessions. For example, previous patient-reported facts can help the Decision Agent route a new question, while current image findings can guide follow-up text reasoning within the same session. However, memory is used as supporting context rather than as definitive clinical evidence.

Overall, the benefit of agent cooperation is that the system can combine *accuracy-oriented components* and *reasoning-oriented components*. Segmentation improves spatial precision, VQA provides language-based interpretation, KG improves structured medical reasoning, RAG provides explanatory grounding, and the Conversation Agent manages missing information. The multi-agent architecture is therefore valuable not only because it is modular, but because it allows specialized agents to produce intermediate evidence that other agents can inspect, reuse, and evaluate.

3.4 Decision Agent

The Decision Agent is responsible for determining which agent should handle a request. For text-only queries, it decides whether the query should be handled by Conversation Agent, KG/RAG reasoning, or Web Search. For image queries, it routes to General Medical Image Agent, Polyp Segmentation Agent, or Polyp VQA Agent.

The decision process considers the input type, the user’s intent, keywords related to segmentation or mask generation, and the context of the conversation. In future extensions, long-term memory can further improve routing by providing patient-specific background before the decision is made.

Routing errors are important because they can create downstream failures. If a symptom consultation is routed directly to answer generation, the system may skip necessary follow-up questions. If a general image is routed to the specialized polyp workflow, the segmentation model may produce meaningless output. If a polyp VQA question is routed to a generic VLM, the answer may lack domain-specific reasoning. Therefore, the Decision Agent is evaluated not only as a convenience component but also as a safety-critical part of the architecture.

The routing strategy uses a conservative principle: when the task is ambiguous, the system should prefer clarification or a safer general path over a highly specialized agent. Specialized agents are useful when the intent and input type match their assumptions. Otherwise, they can

create a false sense of precision.

3.5 Knowledge and Context Sources

The system combines several knowledge and context sources:

- **RAG documents:** unstructured Vietnamese medical information stored in a vector database;
- **Knowledge Graph:** structured disease-symptom-drug relations;
- **Web Search:** fallback source for information not covered by internal data;
- **Short-term state:** conversation context within the current workflow;
- **Long-term memory:** patient-related facts retrieved through Mem0;
- **Visual masks:** segmentation outputs used as grounding signals for Polyp VQA.

This combination allows the system to support both explanatory medical answers and structured reasoning.

The system treats these sources as context providers with different reliability profiles. RAG can provide rich explanations but may retrieve irrelevant text. The Knowledge Graph provides explicit relations but may be incomplete. Web Search can cover recent or missing information but requires careful source filtering. Memory can personalize the response but may contain outdated or user-reported facts. Visual masks provide spatial grounding but depend on segmentation quality. The final response should therefore reflect the strength and uncertainty of each source.

3.6 Safety and Human-in-the-loop Validation

The system includes safety-oriented response behavior. It provides medical disclaimers, recommends professional consultation when necessary, and avoids presenting AI output as a definitive diagnosis. The human-in-the-loop validation mechanism allows users to confirm or comment on generated responses. This feedback can support later quality review and can help determine which facts are reliable enough to be stored in long-term memory.

Human validation is placed after the AI response rather than inside every internal reasoning step. This keeps the system usable while still allowing feedback to be collected. In the scope of this thesis, validation is treated as a safety and review mechanism, not as proof of clinical correctness. The system remains a decision-support tool and should not be interpreted as replacing professional medical judgment.

4 Chapter 4. Text Module: Medical Knowledge Retrieval and Reasoning

4.1 Role and Processing Flow of the Text Module

The Text Module is responsible for medical reasoning over user-provided language. It handles questions about symptoms, diseases, drugs, medical concepts, and follow-up concerns. Its main goal is to produce responses that are grounded, cautious, and aware of missing clinical information.

The general processing flow is:

1. analyze the user query and conversation context;
2. determine whether the query contains enough clinical information;
3. ask follow-up questions if information is insufficient;
4. retrieve relevant knowledge from RAG and/or Knowledge Graph;

5. use Web Search if internal sources are insufficient;
6. synthesize a structured and safe medical response.

This flow is designed to make the Text Module cautious by default. The module does not assume that every query should immediately trigger a final answer. Instead, it separates the tasks of understanding the user intent, checking information completeness, retrieving evidence, and generating the final response.

The Text Module therefore acts as the main reasoning layer for non-image medical queries. It must handle two very different types of user input. The first type is educational, such as asking what a disease is or how a medication works. The second type is consultation-oriented, such as describing symptoms and asking what might be happening. The second type requires stricter information sufficiency checks because the same symptom can have different meanings depending on duration, severity, age, medical history, and associated signs.

The workflow in Figure 2 also shows that text reasoning is not reduced to document retrieval alone. The Text Module combines conversational intake, KG/RAG reasoning, information assessment, web fallback, and final medical response generation. This layered structure is important because the main risk in medical QA is not only whether the system knows a fact, but whether it knows when a fact is insufficient for a patient-specific conclusion.

4.2 Text Query Types and Routing Behavior

The Text Module is designed around the observation that medical text queries are not homogeneous. Different query types require different agent behavior. A general educational question can usually be answered after retrieval, while a patient-specific symptom question may require follow-up questioning before retrieval. A medication-related question may require checking contraindications, allergy context, and source reliability. A non-medical or socially conversational input should not be forced into a medical reasoning pipeline.

Query type	Main agent behavior	Reasoning requirement
Medical education	Use RAG and/or KG to retrieve reliable medical knowledge.	Provide grounded explanation without over-personalizing the answer.
Symptom consultation	Assess information sufficiency and ask follow-up questions if needed.	Avoid premature diagnosis when onset, severity, duration, risk factors, or warning signs are missing.
Structured relation query	Route to KG reasoning and use RAG for explanation when needed.	Connect entities such as symptoms, diseases, drugs, body parts, and risk factors.
Medication or treatment question	Use retrieval cautiously and consider patient context such as allergies or chronic diseases when available.	Avoid unsafe recommendations and encourage professional consultation for dosage or treatment decisions.
Missing or non-medical input	Use the conversational agent or request clarification.	Prevent irrelevant retrieval and keep the system aligned with user intent.

Table 2: Text query types and corresponding routing behavior

Table 2 shows why routing is part of the reasoning problem. If every text query is treated as a retrieval query, the system may return medical documents even when the user actually needs clarification. Conversely, if every symptom query is treated as a conversational query,

the system may underuse available medical evidence. The Text Module therefore uses routing to decide when to ask, when to retrieve, when to reason over graph relations, and when to synthesize a response.

4.3 Conversation Agent and Multi-hop QA

The Conversation Agent handles cases where the user’s query is incomplete or conversational. In medical consultation, missing information is common. For example, a symptom description may lack onset time, severity, location, duration, associated symptoms, medical history, or current medication use.

Multi-hop QA is used to collect this missing information step by step. Instead of generating a diagnostic suggestion immediately, the agent asks focused follow-up questions. This process helps the system approximate a structured clinical intake conversation. The goal is not to replace a doctor, but to reduce the risk of giving an answer from insufficient evidence.

The follow-up questions are designed to be concise and clinically relevant. For symptom-based consultation, the system prioritizes red flags and high-impact missing variables. For example, in a fever-related query, duration, peak temperature, associated respiratory or gastrointestinal symptoms, medication use, and risk factors may change the response. In a pain-related query, location, severity, onset, radiation, and accompanying symptoms are important.

This behavior improves usability because the system does not overwhelm the user with a full medical form. It asks only the next useful question based on the current context.

The multi-hop design also improves traceability. Each follow-up question corresponds to a missing clinical variable. This makes it possible to analyze whether the agent is asking useful questions or merely extending the conversation. In evaluation, this behavior is reflected by Multi-Hop Awareness Score, which rewards the system for recognizing when additional information is needed.

4.3.1 Illustrative Multi-hop Consultation Flow

An illustrative symptom consultation demonstrates why multi-hop behavior is needed. Suppose the user writes: “I have stomach pain and nausea. What disease could it be?” A single-turn chatbot may immediately list gastritis, food poisoning, reflux, or appendicitis. However, this answer is clinically weak because the initial input lacks duration, pain location, severity, fever, vomiting, stool changes, medication use, pregnancy status, and warning signs.

In the proposed workflow, the Conversation Agent first treats the query as underspecified. Instead of giving a strong conclusion, it asks a focused follow-up question such as: “Where is the pain located, how long has it lasted, and do you have fever, vomiting blood, black stool, severe pain, or persistent vomiting?” This question is not arbitrary; it targets variables that can change the safety framing of the response. If the user later reports mild epigastric pain after meals without red flags, the system can retrieve information about gastritis or reflux and provide cautious advice. If the user reports severe right-lower abdominal pain with fever, the system should prioritize urgent professional care rather than ordinary health education.

This example illustrates the difference between answer generation and consultation support. The goal of the Conversation Agent is not to prolong the dialogue, but to determine whether the current evidence is sufficient for the next safe action. The final response may still include possible causes, but it should clearly separate likely explanations, missing information, warning signs, and next steps.

4.4 Information Sufficiency Assessment

Information sufficiency is the mechanism that decides whether the system has enough context to answer. If the available information is not enough, the system should ask for additional details. If the information is sufficient, it can proceed to retrieval and response generation.

This mechanism is important because many medical questions should not be answered directly. A safe assistant should distinguish between:

- general medical education questions that can be answered from knowledge sources;
- symptom-based questions that require follow-up details;
- high-risk cases that should recommend immediate professional care;
- out-of-scope or non-medical queries.

Information sufficiency is also connected to evaluation. A model that refuses to conclude when the evidence is incomplete may appear less decisive, but it is safer in medical reasoning. Therefore, this thesis uses Multi-Hop Awareness Score to measure whether the model can recognize insufficient information rather than forcing an answer.

In implementation, information sufficiency should be interpreted as a decision boundary rather than a fixed checklist. Some questions can be answered with little personal context because they are general medical education questions. Other questions require more detail because they imply patient-specific advice. For example, explaining what gastritis is does not require a full intake dialogue, but advising a user with severe abdominal pain requires caution, red-flag screening, and possibly urgent-care guidance.

An important insight is that information sufficiency is an architectural control point, not only a prompt instruction. In the workflow, the result of this assessment changes the route of execution: the system may continue to MedGemma-based synthesis, ask the conversational agent to collect missing details, or invoke Web Search when internal evidence is weak. This makes the Text Module more similar to a clinical triage workflow than to a single-turn QA model.

The information sufficiency step also prevents a common failure mode of retrieval-augmented systems. A retrieved passage can make a response appear grounded even when the original user question is incomplete. For example, if the user says only “I have chest pain,” RAG may retrieve documents about angina, reflux, anxiety, or muscle pain. These documents are medically relevant, but they do not solve the missing-context problem. The system must first identify urgent warning signs and ask about onset, severity, radiation, shortness of breath, sweating, and cardiovascular history. Therefore, evidence availability and clinical sufficiency are related but not identical.

In this thesis, sufficiency is treated as a dynamic threshold. The threshold is lower for general education and higher for patient-specific recommendations. It is also higher when the symptom category is associated with potentially urgent conditions. This design makes the Text Module cautious without preventing it from answering low-risk factual questions.

4.5 Medical Knowledge Retrieval with RAG and Knowledge Graph

The system uses both RAG and Knowledge Graph retrieval because they provide complementary strengths.




Figure 3: Text Module architecture for medical knowledge retrieval




Figure 4: Parallel KG-RAG processing in the Text Module

The parallel design in Figure 4 reflects a central assumption of this thesis: unstructured

and structured medical knowledge are complementary. RAG is strong when the system needs explanatory evidence, while KG reasoning is strong when the question depends on explicit relations between entities such as symptoms, diseases, drugs, and body parts. Running these branches in parallel allows the system to compare different evidence forms before generating an answer.

RAG is effective when the answer requires explanatory medical text. It retrieves semantically relevant documents from the vector store and provides them as context for response generation. This is useful for disease descriptions, medication explanations, and general medical education.

Knowledge Graph retrieval is effective when the query involves structured relationships. For example, questions about which symptoms are associated with a disease, or which diseases may be linked to a set of symptoms, can benefit from explicit graph relations.

Query expansion and context filtering are used to improve retrieval quality. Query expansion adds related medical terms and reformulates the query, while context filtering reduces irrelevant or noisy retrieved results.

The RAG and KG components also support different explanation styles. RAG is better suited for patient-friendly paragraphs because it retrieves natural-language medical content. KG is better suited for compact relational evidence because it retrieves structured connections. A robust medical answer can use both: graph relations for reasoning and retrieved documents for explanation.

Query expansion is especially useful for Vietnamese medical text because users may describe symptoms in informal language while the dataset uses more formal medical terminology. Expansion can add synonyms, related symptoms, disease names, or anatomical terms. For example, a user may describe “đau bao tử” while a medical corpus may contain entries related to stomach pain, epigastric pain, gastritis, or gastroesophageal reflux. Expansion improves the chance that retrieval reaches the relevant region of the knowledge base.

Context filtering is the opposite operation: it removes retrieved passages that are semantically close but clinically irrelevant. In medical QA, this is necessary because two passages can share vocabulary while discussing different conditions, populations, or levels of severity. Filtering also reduces prompt noise, which helps the response generator focus on evidence that is actually useful.

The evaluation in Chapter 8 supports the need for this hybrid design. KG + MedLLM achieves the strongest result on both One-Hop Accuracy and Multi-Hop Awareness Score, suggesting that explicit graph relations are valuable for structured medical reasoning. At the same time, RAG remains necessary because graph triples alone rarely provide enough explanatory detail for patient-facing responses. Therefore, the contribution of the Text Module is not simply choosing KG over RAG, but coordinating them according to the evidence need of each query.

4.5.1 Coordinated KG-RAG Evidence Construction

The KG and RAG branches contribute different forms of evidence. The KG branch is useful for identifying candidate relations, such as a disease being associated with a symptom or a drug being related to a disease category. The RAG branch is useful for converting those relations into explanatory language that can be understood by non-expert users. This distinction matters because a medical assistant must both reason and communicate.

For example, when the user asks about a symptom cluster, the KG branch can return candidate diseases connected to the symptoms. These candidates should not be treated as diagnoses. Instead, they serve as structured hypotheses that guide retrieval and response generation. The RAG branch can then retrieve documents explaining the candidate conditions, typical symp-

toms, warning signs, and general management. The response generator can synthesize these sources into a cautious answer: possible explanations, why they may be relevant, what information is missing, and when to seek medical care.

This coordination also reduces two opposite risks. If the system relies only on KG, the answer may become too sparse and relation-like, lacking patient-friendly explanation. If the system relies only on RAG, the answer may become verbose and less structured, especially when retrieved passages contain overlapping or weakly relevant information. By combining both, the Text Module can use KG for structure and RAG for explanation.

Stage	Agent output	Role in final response
Query Expansion	Medical synonyms, related symptoms, disease names, and anatomical terms.	Bridges colloquial user language and formal medical terminology.
KG Agent	Structured entity relations and candidate disease-symptom links.	Provides compact reasoning structure and candidate hypotheses.
RAG Agent	Retrieved explanatory passages from medical documents.	Provides patient-friendly explanation and supporting context.
Information Assessment	Decision on whether evidence is sufficient, weak, conflicting, or missing.	Determines whether to answer, ask follow-up questions, or use fallback search.
Response Generator	Structured medical response with uncertainty and safety framing.	Converts evidence into an understandable answer without presenting it as a definitive diagnosis.

Table 3: Coordinated evidence construction in the Text Module

4.6 Parallel Reasoning and Fallback Strategy

The system can evaluate RAG and Knowledge Graph results in parallel. If one source provides sufficient information, it may be selected as the main evidence source. If internal sources are insufficient, the system can use Web Search as a fallback.

This strategy reduces dependence on any single source. RAG may fail if the relevant document is not retrieved. The Knowledge Graph may fail if the needed relation is missing. Web Search may provide updated information but requires careful filtering. By combining these sources, the Text Module becomes more robust.

The fallback strategy is especially important for medical topics that change over time or are not well represented in the internal data. However, fallback search should be treated cautiously. The system should prefer trusted medical sources and avoid giving strong claims when source quality is uncertain.

Parallel reasoning does not mean that all sources are merged blindly. The system should compare source relevance, consistency, and sufficiency. When KG and RAG agree, the answer can be more confident. When they conflict or when one source is weak, the response should explain uncertainty or ask for more information. This is important because adding more sources can increase coverage but can also introduce contradictions.

This strategy also gives the system an interpretable failure path. If the KG branch returns no relevant relation, the system can still use RAG for broader explanatory context. If RAG retrieves weak or overly general passages, the KG branch may provide a more precise disease-symptom relation. If both internal sources are insufficient, the Web Search Agent becomes a

controlled fallback rather than an uncontrolled first response. As a result, robustness comes from source diversity and routing logic, not from assuming that one knowledge source is universally reliable.




Figure 5: Medical chain-of-thought and information sufficiency flow

4.7 Safe Medical Response Generation

The final response is generated from the available evidence and conversation context. A safe medical response should:

- answer the user’s question directly when possible;
- explain the reasoning in understandable language;
- cite or mention sources when available;
- state uncertainty when information is incomplete;
- recommend professional care for severe or urgent symptoms;
- avoid presenting the AI output as a definitive diagnosis.

The generated response is then available for human-in-the-loop validation.

The final response is therefore the result of a controlled reasoning pipeline rather than a raw model completion. This design helps separate evidence collection from natural-language generation, making the system easier to analyze and improve.

The response format should be adapted to risk level. For low-risk educational questions, the response can provide a concise explanation and general advice. For symptom-based questions, the response should include what is known, what is missing, possible warning signs, and when to seek professional care. For urgent symptoms, the system should prioritize escalation over detailed explanation.

The response generator should also preserve the provenance of reasoning as much as possible. If the answer is mainly based on RAG evidence, it should emphasize retrieved medical information. If the answer depends on KG relations, it should explain that the listed conditions are related possibilities rather than confirmed diagnoses. If Web Search is used, the response should be more conservative because source quality may vary. This provenance-aware response

style improves user trust and makes the system easier to audit.

In addition, the response should separate *medical knowledge* from *patient-specific interpretation*. For example, the system can explain that gastritis commonly causes epigastric pain and nausea, but it should not conclude that the user has gastritis without sufficient context. This separation is central to safe medical assistance: the assistant may provide useful education and triage guidance, but it should avoid replacing clinical diagnosis.

5 Chapter 5. Vision Module: General Medical Image Analysis, Polyp Segmentation, and Polyp VQA

5.1 Role and Processing Flow of the Vision Module

The Vision Module processes image-based medical inputs. It supports general medical image interpretation and a specialized colonoscopy polyp workflow. The module is designed to coordinate visual models with the multi-agent architecture so that image findings can be combined with text questions and patient context.

The module follows a task-specific design. General medical image interpretation is kept broad and cautious, while polyp-related processing is more specialized and more formally evaluated. This allows the system to support multiple image inputs without treating all image types as equally reliable.

The design intentionally separates general image analysis from specialized polyp analysis. A general VLM may be useful for describing visible patterns, but colonoscopy polyp segmentation and Polyp VQA require domain-specific assumptions, datasets, and evaluation metrics. This separation makes the system more transparent: the user and evaluator can identify whether an answer comes from a general visual model or from the specialized polyp workflow.




Figure 6: Vision Module workflow for medical image analysis

5.2 Image Routing and Visual Task Classification

When an image is uploaded, the system determines whether the task is general image interpretation, polyp segmentation, or Polyp VQA. If the user explicitly requests mask generation or segmentation, the image is routed to the Polyp Segmentation Agent. If the user asks a question about an endoscopic polyp image, the image is routed to the Polyp VQA Agent. Other medical images can be handled by the General Medical Image Agent.

This routing step is important because the same image can support different tasks. A user may ask for a segmentation mask, a description of visible findings, or a question about risk. Routing prevents the system from applying the wrong agent to the wrong task.

Visual routing also depends on the accompanying text. An image without a question may require general description, while the same image with a question about polyp type should trigger Polyp VQA. If the user asks specifically for localization or mask generation, segmentation becomes the primary task. The image and the text prompt are therefore interpreted together rather than independently.

5.3 General Medical Image Agent

The General Medical Image Agent handles medical images outside the specialized polyp workflow. It uses a multimodal model to describe visible findings, provide cautious interpretation, and recommend appropriate follow-up when necessary.

This agent is positioned as a support tool rather than a diagnostic authority. Its output

should be interpreted as preliminary information and should not replace clinical review by trained professionals.

Because the General Medical Image Agent covers a broad space of image modalities, it is not the main formal evaluation target in this thesis. Its role is to provide system coverage and demonstrate extensibility of the vision pipeline. The specialized polyp workflow receives more detailed quantitative evaluation.

The agent should avoid unsupported claims about diagnosis, severity, or treatment. Its most appropriate output style is a cautious description of visible content, possible medical relevance, and recommendation for professional interpretation when the image is clinically significant. This safety framing is necessary because different medical image modalities require different expert knowledge and acquisition conditions.

5.4 Polyp Segmentation Agent

The Polyp Segmentation Agent is designed for colonoscopy images. Its task is to segment polyp regions and distinguish between background, neoplastic polyp regions, and non-neoplastic polyp regions. The segmentation output is useful for two purposes: direct visual interpretation and visual grounding for Polyp VQA.

This agent converts an image-level input into a spatially localized representation. That representation is easier to inspect than a text-only answer and can be passed to other agents. In the proposed architecture, segmentation is therefore both a prediction task and an intermediate reasoning artifact.

The three-class design is also useful for downstream reasoning. A binary mask can indicate where a polyp is located, but it does not distinguish neoplastic from non-neoplastic regions. By preserving class information, the mask can provide richer evidence to the Polyp VQA Agent. This is particularly relevant when the user’s question asks about suspicious areas or class-related interpretation.

5.5 DeepLabV3+ for Polyp Segmentation

DeepLabV3+ is used as the core segmentation architecture for the Polyp Segmentation Agent. The target output is a dense three-class mask in which each pixel is assigned to background, neoplastic polyp region, or non-neoplastic polyp region. This setting requires both accurate localization and stable class prediction inside the lesion area.

Architecture objective. DeepLabV3+ is selected because it combines high-level semantic context with boundary recovery. Colonoscopy frames often contain lesions with variable size, weak contrast, specular highlights, and unclear tissue boundaries. A useful architecture must therefore capture multi-scale context while still preserving enough spatial detail for pixel-level prediction.

Main components. The architecture contains four important parts:

- **Backbone encoder:** extracts hierarchical feature maps from the input colonoscopy image. The deeper layers provide semantic information, while lower-level features retain edge and texture details.
- **Atrous convolution:** expands the receptive field without excessively reducing feature-map resolution. This is useful when a polyp boundary requires both local texture cues and wider mucosal context.
- **ASPP:** Atrous Spatial Pyramid Pooling applies parallel atrous convolution branches at multiple dilation rates, together with image-level pooling, to capture objects and contextual patterns at different scales.
- **Decoder:** combines upsampled high-level features with low-level features to recover spa-

tial detail and produce sharper segmentation boundaries.

The segmentation agent follows the DeepLabV3+ encoder-decoder design, summarized in Figure 7. The encoder first extracts hierarchical features from the input image through a deep convolutional backbone with atrous convolution. Atrous convolution is important because it enlarges the effective receptive field without reducing the feature-map resolution too aggressively. In colonoscopy frames, this is useful because the model must capture both local texture changes around the lesion boundary and broader contextual information from the surrounding mucosa.




Figure 7: DeepLabV3+ encoder-decoder architecture with Atrous Spatial Pyramid Pooling

Processing workflow. Let X be the input colonoscopy image. The encoder first produces high-level semantic features:

$$F_e = ASPP(DCNN(X)). \quad (1)$$

The ASPP output is then upsampled and concatenated with low-level features F_l extracted from earlier layers:

$$F_d = \text{Concat}(\text{Upsample}(F_e), \text{Conv}_{1 \times 1}(F_l)). \quad (2)$$

The decoder refines the combined representation and produces the final pixel-wise prediction:

$$P = \text{Upsample}(\text{Conv}_{3 \times 3}(F_d)). \quad (3)$$

The dense prediction P is converted into the color mask used by the rest of the system. Blue denotes background, red denotes neoplastic regions, and green denotes non-neoplastic regions. This color convention is not only for visualization; it also provides the symbolic input used by the post-processing method and the segmentation-grounded Polyp VQA setting.

Relevance to the thesis. The encoder-ASPP-decoder design matches the needs of three-class polyp segmentation. ASPP helps detect lesions across scale variation, the decoder helps recover boundaries, and the class-wise mask preserves information needed for neoplasm characterization. However, a dense neural prediction can still contain small inconsistent class fragments inside a connected lesion region. The post-processing method below is therefore introduced as a region-level refinement step after DeepLabV3+ prediction.

5.6 Post-processing and Overlay Mask

Raw segmentation predictions may contain noisy regions or inconsistent labels inside a connected polyp region. To improve consistency, the system applies post-processing based on two main steps.

Connected Component Analysis identifies non-background connected regions in the mask. Each connected component is treated as a candidate object region. Let $M(x, y)$ denote the raw predicted color label at pixel (x, y) , where blue represents background, red represents neoplastic tissue, and green represents non-neoplastic tissue. The first step converts the predicted mask into a binary region-of-interest map:

$$I(x, y) = \begin{cases} 1, & \text{if pixel } (x, y) \text{ belongs to a target region (non-blue),} \\ 0, & \text{otherwise.} \end{cases} \quad (4)$$

Two pixels $p = (x_1, y_1)$ and $q = (x_2, y_2)$ are considered connected under an 8-neighborhood rule if:

$$|x_1 - x_2| \leq 1, \quad |y_1 - y_2| \leq 1, \quad I(p) = I(q) = 1. \quad (5)$$

The connected foreground pixels form a set of connected components:

$$\{R_i\}_{i=1}^N \quad \text{such that} \quad \forall (p, q) \in R_i, \quad p \leftrightarrow q, \quad (6)$$

where R_i is the i -th connected region, N is the number of detected connected regions, and $p \leftrightarrow q$ means that there exists a path of connected foreground pixels between p and q . The region label map is then defined as:

$$L(x, y) = i \quad \text{if} \quad (x, y) \in R_i. \quad (7)$$

Region-wise Color Dominance Refinement checks the number of red and green pixels inside each connected component. For each connected region R_i , the number of pixels assigned to class k is computed as:

$$C_k(R_i) = \sum_{(x,y) \in R_i} \mathbf{1}[M(x, y) = k], \quad (8)$$

where $k \in \{\text{red}, \text{green}\}$ and $\mathbf{1}[\cdot]$ is an indicator function that returns 1 when the condition is true and 0 otherwise. The dominant class of the region is:

$$k^* = \arg \max_k C_k(R_i). \quad (9)$$

The same counts also provide a simple measure of how reliable the region-level decision is. Let $|R_i|$ be the number of pixels in region R_i . The dominance ratio is:

$$\rho_i = \frac{\max_k C_k(R_i)}{|R_i|}. \quad (10)$$

A high ρ_i means that most pixels in the connected region already agree on the same class, so refinement mainly removes small noisy fragments. A low ρ_i indicates that the region is ambiguous because red and green occupy comparable portions. In that case, assigning a single class is more risky and should be interpreted carefully in qualitative analysis.

The refined mask M' is obtained by assigning the dominant class to all pixels in the connected component:

$$M'(x, y) = k^*, \quad \forall (x, y) \in R_i. \quad (11)$$

These equations make the assumption of the method explicit. Connected component analysis defines the spatial unit of refinement through R_i , while color dominance refinement defines the semantic decision through $C_k(R_i)$ and k^* . Therefore, the method does not modify the mask arbitrarily at pixel level. It first asks whether pixels belong to the same foreground object, then assigns that object the class that already dominates inside the raw prediction. When ρ_i is high, isolated minority-class pixels are treated as local noise. When ρ_i is low, the processed mask may hide a genuinely uncertain prediction, so those cases should be inspected as possible failure cases.

The formulation provides three practical insights for interpreting the post-processing result:

- **Spatial insight from R_i :** refinement is performed per connected foreground object, not across the whole image. Therefore, separate candidate lesions are processed independently.
- **Semantic insight from $C_k(R_i)$:** the corrected class is determined by the class distribution already present in the raw prediction. The method regularizes the model output but does not introduce external clinical knowledge.
- **Reliability insight from ρ_i :** high-dominance regions are good candidates for refinement, while low-dominance regions should be treated as uncertain cases where the raw prediction itself is class-ambiguous.

Representative post-processing behavior is provided in Figure 8. Each row contains the original colonoscopy image, the ground-truth mask, the raw model prediction, and the processed prediction. The visualization makes the purpose of the method clear: the post-processing step does not attempt to redraw the lesion from scratch; instead, it regularizes the semantic label inside each connected predicted region.




Figure 5: Comparison Result of The Masking Segmentation with The Technique

Figure 8: Visualization of raw and post-processed polyp segmentation masks

In the first and second rows, the raw prediction contains a connected region with both red and green pixels. According to the formulation above, these pixels belong to the same R_i , and the dominant count $C_k(R_i)$ determines whether the region becomes red or green after refinement. After color dominance refinement, the processed prediction assigns the dominant class to the whole connected component, reducing the mixed-label artifact. In the third row, the raw prediction contains a small incorrect class fragment attached to the main object. Because this fragment is a minority label inside the same connected component, it is removed while the overall shape is preserved. In the fourth row, the method keeps multiple connected lesion-like structures but corrects local class noise inside the larger connected area.

These examples show that post-processing is most effective when the segmentation model has already localized the lesion reasonably well but produces unstable class labels within the region. The technique is therefore a region-consistency correction step rather than a detector. It can improve the readability of the mask and the Dice Score, but it remains dependent on the quality of the raw prediction.

After refinement, the mask is overlaid on the original image. The overlay helps users inspect the model output and provides a grounded visual input for the Polyp VQA Agent.

The overlay step is particularly useful for interpretability. Instead of returning only a numerical prediction, the system provides a visual explanation of where the model focuses. This is important in a medical context because users and reviewers need to inspect whether the

model’s attention is plausible.

Post-processing also helps align the segmentation output with the assumptions of visual grounding. A VQA model that receives a fragmented or mixed-color overlay may be guided toward inconsistent evidence. A cleaner overlay makes the grounded input easier for both humans and models to interpret. However, this method is intentionally conservative: it improves regional consistency but does not create new lesion regions. If the initial model misses the polyp entirely or detects a false foreground region, connected component analysis cannot correct that semantic error. Qualitative examples should therefore still show failure cases.

5.7 Polyp VQA Agent Design with LoRA Fine-tuned VLMs

The Polyp VQA Agent answers natural-language questions about colonoscopy polyp images. It is invoked when the Decision Agent detects that the user is asking a question about an endoscopic polyp image rather than requesting only segmentation or general image description. The agent receives an image-question input, formats it for the selected VLM endpoint, and returns a concise answer that should be medically relevant and visually grounded in the supplied image.

The agent uses LoRA fine-tuned VLMs hosted through a configurable base URL. In this thesis, the VQA fine-tuning is designed for two 4B-scale backbones: MedGemma-4B-it and Qwen3.5-4B. This design allows the system to compare a medically oriented VLM backbone with a strong general instruction-following backbone while keeping the agent interface compatible with the rest of the workflow.

The dataset construction and experimental split for Polyp VQA are presented in Chapter 7. In this chapter, the focus is the agent role: converting image-question inputs into clinically relevant answers and integrating the fine-tuned VLM with the segmentation workflow.

The VQA agent is connected to the segmentation workflow through visual grounding. When a segmentation mask is available, the agent can receive both the original image and the overlay. This design gives the VLM explicit visual evidence about the region of interest while preserving the original image as the primary visual context. If no segmentation mask is available, the agent falls back to the original-image VQA setting.

This design gives the polyp workflow a two-stage reasoning structure. The Segmentation Agent contributes accuracy at the visual-localization level by identifying where the relevant lesion-like region is located. The Polyp VQA Agent contributes reasoning at the semantic level by answering clinical questions about the image and the highlighted region. The final behavior is therefore stronger than either component alone: segmentation without VQA is difficult for users to interpret in natural language, while VQA without segmentation may rely on less explicit visual evidence.

Using hosted fine-tuned VLMs through a base URL keeps the agent workflow independent from the physical deployment of the models. From the agent’s perspective, each Polyp VQA model is a callable service that receives an image-question input and returns an answer. The thesis focuses on the AI behavior of these models and their integration with segmentation, not on web backend details.

5.8 Segmentation-to-VQA Agent Cooperation

The cooperation between the Polyp Segmentation Agent and the Polyp VQA Agent is one of the clearest examples of why a multi-agent architecture is useful. The two agents operate on the same image but solve different subproblems. The Segmentation Agent answers a localization question: which pixels correspond to relevant polyp regions, and which predicted class is assigned to those regions? The VQA Agent answers a semantic question: how should the visual

evidence be described in response to a natural-language question?

This separation produces a useful intermediate evidence object: the segmentation overlay. The overlay has three roles. First, it provides spatial grounding by highlighting the region that the downstream model should inspect. Second, it improves interpretability because users and reviewers can see whether the model is focusing on a plausible area. Third, it creates a bridge between pixel-level prediction and language-level reasoning. The VQA Agent does not need to infer the region of interest only from global image features; it receives an explicit visual cue generated by the segmentation pipeline.

Agent	Main contribution	Limitation when used alone
Polyp Segmentation Agent	Produces pixel-level localization and class-colored masks for neoplastic or non-neoplastic regions.	A mask alone does not explain findings, answer user questions, or provide natural-language clinical reasoning.
Polyp VQA Agent	Converts image-question input into a clinically relevant natural-language answer.	Without grounding, the VLM may rely on global context, dataset priors, or visually irrelevant regions.
Combined workflow	Uses the segmentation mask as visual evidence for VQA reasoning.	The workflow depends on mask quality; an incorrect mask can mislead the VQA response.

Table 4: Complementary roles of the Polyp Segmentation Agent and Polyp VQA Agent

Table 4 summarizes the complementarity between the two agents. The Segmentation Agent improves the precision of visual localization, while the VQA Agent improves the interpretability of the output. Their cooperation is especially relevant for questions such as whether a highlighted region is suspicious, where the polyp is located, or whether the visible region appears consistent with a neoplastic or non-neoplastic class label. These questions require both spatial evidence and language reasoning.

The cooperation also changes how the system should be evaluated. If segmentation and VQA are evaluated separately, the system may appear strong even when the connection between the two agents is weak. A segmentation model may achieve a high Dice Score but still produce overlays that are confusing for VQA. A VQA model may produce fluent answers but fail to use the highlighted region. Therefore, this thesis evaluates not only the segmentation model and VQA model independently, but also the visual-grounded VQA setting where the mask becomes part of the VQA input.

The experimental result in Chapter 8 supports this design choice. On the mask-matched subset, adding visual grounding improves the LLM Judge score compared with the LoRA-only VQA baseline. This suggests that the segmentation mask provides useful spatial evidence when it is available and aligned with the question. The improvement is strongest in factual accuracy and completeness, which is consistent with the hypothesis that grounding helps the model select relevant visual information rather than merely changing the response style.

However, this cooperation is conditional rather than automatically beneficial. If the Segmentation Agent misses the lesion, highlights irrelevant mucosa, or assigns an incorrect class to the region, the VQA Agent may produce a more confident but wrong answer. This is a typical risk of pipeline-based reasoning: an upstream error can propagate downstream. For this reason, the segmentation overlay should be interpreted as an evidence cue, not as unquestionable ground truth. The VQA response should remain cautious, especially when the mask is small, fragmented, uncertain, or inconsistent with the visible image.

5.8.1 Example Interaction Pattern

An example workflow illustrates the interaction between the two agents. First, the user uploads a colonoscopy image and asks whether there is a visible polyp. The Decision Agent routes the request to the polyp workflow. The Segmentation Agent generates a mask and overlay, which identifies the suspected region. The Polyp VQA Agent then receives the original image, the overlay, and the user’s question. Instead of answering only from the raw image, it can refer to the highlighted region and provide a response such as: the image contains a region that should be reviewed, the highlighted area may correspond to a polyp-like structure, and a clinician should confirm the finding.

The user may then ask a follow-up question: “Why is this region suspicious?” In this case, the system can reuse the same image-derived context. The follow-up question is textual, but it still depends on visual evidence. The workflow can route the question back to the Polyp VQA Agent or combine the VQA output with the Text Module for explanatory medical information. This illustrates cross-turn and cross-modal reasoning: the image is not treated as a one-time input, but as part of the ongoing consultation context.

This interaction pattern is useful for user-facing medical assistance because patients rarely ask only one complete question. They may upload an image, ask what it shows, ask whether the highlighted region is dangerous, ask what symptoms should be monitored, and ask when to see a doctor. The multi-agent workflow supports this behavior by preserving intermediate evidence and allowing later questions to reuse it.

5.9 Original-image and Segmentation-grounded VQA Settings

The Polyp VQA model is studied under two input settings.

Original image + question. The model receives only the original colonoscopy image and the user’s question. This represents a VQA setting where the model must identify relevant visual evidence by itself.

Original image + segmentation mask + question. The model receives both the original image and the segmentation overlay. The mask acts as a visual grounding signal that guides the model toward the polyp region. This setting represents cooperation between the Polyp Segmentation Agent and the Polyp VQA Agent.

Visual grounding can improve answer quality when the mask correctly highlights clinically relevant regions. However, it can also introduce risk if the segmentation mask is inaccurate. Therefore, both settings should be evaluated and compared.

The comparison between these two settings is one of the central experiments in the Vision Module. It tests whether agent cooperation improves multimodal reasoning, rather than assuming that adding more visual inputs always helps.

The second setting represents agent cooperation because the VQA Agent does not create the grounding signal itself. It receives a spatial cue generated by the Segmentation Agent. If the grounded setting improves answer quality, this suggests that modular visual reasoning can be beneficial. If it does not improve or degrades performance, the analysis can identify whether the issue comes from segmentation errors, overlay design, or VLM sensitivity to additional visual inputs.

This separation also makes the evaluation more interpretable. The original-image setting measures the VLM’s baseline ability to extract relevant evidence from the endoscopic frame. The segmentation-grounded setting measures the effect of adding explicit spatial evidence. Therefore, the comparison isolates visual grounding as an intervention in the VQA pipeline rather than mixing it with unrelated changes in model interface or agent orchestration.

6 Chapter 6. Memory Layer: Long-Term Context for Medical Agents

6.1 Role of the Memory Layer

The Memory Layer provides long-term context for the agent system. It stores clinically useful facts that may remain relevant across multiple sessions, such as reported symptoms, allergies, medication use, chronic diseases, treatment reactions, warning signs, and communication preferences.

Mem0 is used as the memory management layer. It does not replace LangGraph, RAG, the Knowledge Graph, or patient data storage. Instead, it provides additional context that can improve personalization and continuity.

A key design principle is separation of responsibility. LangGraph manages workflow state and agent coordination. RAG and the Knowledge Graph provide medical knowledge. The memory layer stores reusable patient-context facts that are useful for future interactions. This separation prevents long-term memory from becoming an uncontrolled knowledge base or a hidden diagnostic record.

6.2 Integration Architecture

LangGraph remains the central workflow orchestrator. Mem0 is integrated as a supporting memory service that can be called by nodes in the graph. Before routing, the Decision Agent can retrieve memory associated with a patient identifier. Before response generation, relevant memory can be passed into Conversation Agent, KG/RAG reasoning, or vision-related agents as context.

After response generation and validation, new clinically useful information can be considered for memory storage. This creates a feedback loop where long-term context improves future interactions.

The integration follows a retrieval-before-reasoning pattern. At the beginning of a request, the system retrieves relevant memory by patient identifier and semantic similarity. The retrieved memory is then supplied as supporting context to the Decision Agent or task-specific agents. At the end of the workflow, candidate facts are extracted from the interaction and stored only when they are durable, clinically useful, and appropriate for future reuse.

6.3 Technical Pipeline

The memory pipeline consists of fact extraction, embedding, vector storage, and semantic retrieval. An LLM extracts durable clinical facts from conversation content. An embedding model encodes the memory text. A vector store supports semantic search over stored memories.

The quality of memory depends on the extraction LLM, the embedding model, and the custom instructions that define what should or should not be stored. For Vietnamese interactions, the system should preserve patient-reported facts in Vietnamese when appropriate.

This means that Mem0 should be evaluated as an integration component rather than as a new memory algorithm. The system quality depends on how well facts are extracted, how accurately embeddings retrieve related memories, and how strictly memory instructions prevent irrelevant or unsafe content from being stored. In a medical setting, memory policies are as important as retrieval accuracy.

6.4 Impact on Agent Reasoning

Long-term memory can improve several parts of the system.

Decision Agent. Memory can improve routing. For example, if the memory indicates a drug allergy, a medication-related question should be handled with stronger medical reasoning rather than simple conversation.

Conversation Agent. Memory can reduce repeated questions by reusing previously reported facts.

KG/RAG Agents. Memory can personalize query expansion and retrieved context.

Vision-related agents. Memory can provide patient context, but image evidence should remain the primary basis for image interpretation.

For example, if previous sessions indicate that a patient has a history of colon polyps, this context may help the Decision Agent route a new colonoscopy question to the specialized polyp workflow. However, the segmentation mask and the current image should remain the primary evidence for any image-related response. Memory should guide context, not override visual evidence.

6.5 Safety Constraints and Limitations

Memory is not a diagnosis. Patient-reported facts should be represented as reported information, such as “the user reports...” or “the patient states...”. AI-inferred conclusions should be marked as uncertain unless validated.

The memory layer should not store raw images, base64 image content, API keys, phone numbers, addresses, or unnecessary personally identifiable information. Memory can also become outdated or incorrect, so it should be used as supporting context rather than authoritative truth.

Future deployment would require stronger governance mechanisms, including memory review, correction, deletion, source tracking, and confidence marking. These mechanisms are especially important when memory is used across multiple sessions, because a fact that was once correct may become outdated after medication changes, new diagnoses, or updated patient reports.

7 Chapter 7. Data and Experimental Setup

7.1 VimedAQA Dataset

VimedAQA is used as the Vietnamese medical question-answering source for the RAG component [6]. In the local index used by this thesis, the dataset is represented by 44,228 JSONL records across six groups: body part, disease, drug, and medicine-related passages. Each record contains a medical context, title, keyword, topic identifier, source URL, and author metadata. This structure makes the dataset suitable for retrieval-oriented medical QA rather than direct model fine-tuning.




Figure 9: Medical QA dataset overview used for retrieval-oriented experiments

VimedAQA Group	Number of Records
Body part	4,956
Disease 1	7,500
Disease 2	8,154
Drug	9,754
Medicine 1	7,500
Medicine 2	6,364
Total	44,228

Table 5: Local VimedAQA record distribution used for RAG indexing

For retrieval-based reasoning, each record is converted into a retrievable passage. During inference, the user query is embedded and compared with stored passages, and the most relevant contexts are supplied to the answer generator. This setup allows the system to reuse Vietnamese medical information while keeping response generation separate from data indexing.

The quality of this component depends on preprocessing and retrieval design. Noisy passages, duplicated records, or overly broad chunks can reduce retrieval precision. Therefore, the dataset is treated as a knowledge source that requires cleaning, normalization, consistent chunking, and metadata-aware indexing.

7.2 VietMedKG

VietMedKG is used as the structured medical knowledge source for the Knowledge Graph Agent. In the local data used by this thesis, the graph source table contains 8,900 disease records with fields such as disease name, description, disease type, prevention, causes, symptoms, risk groups, comorbidities, treatment methods, clinical department, tests, nutrition advice, and medication-related information. These fields are converted into structured facts for graph-based retrieval and reasoning.




Figure 10: Overview of the VietMedKG data source used by the Knowledge Graph Agent

Unlike RAG, which retrieves semantically similar text, VietMedKG supports relation-oriented reasoning. For example, the system can retrieve symptoms associated with a disease or identify diseases related to a symptom set. This makes the graph particularly useful for one-hop medical relation questions and for checking whether a generated answer is consistent with structured medical facts.

7.3 BKAI-IGH NeoPolyp-Small Dataset

BKAI-IGH NeoPolyp-Small is used for the polyp segmentation experiments. The dataset contains colonoscopy images with pixel-level annotations for neoplastic and non-neoplastic polyp regions. In this thesis, the segmentation task is formulated as a three-class problem: background, neoplastic polyp region, and non-neoplastic polyp region. This makes the dataset suitable for evaluating both lesion localization and class-aware mask prediction.

BKAI-IGH NeoPolyp

A dataset for colonoscopy polyp segmentation and neoplasm characterization

Overview

Data

Code

Models

Discussion

Leaderboard

Rules

Team

Submissions

Dataset Description

File description:

- test:** test images
- train:** training images
- train_gt:** ground-truths of training images
- sample_submission.csv:** sample submission

Ground-truth explanation

- Red color: neoplastic polyps
- Green color: non-neoplastic polyps
- Black color: background

Files

2201 files

Size

375.42 MB

Type

jpeg, csv

License

Subject to Competition Rules

Figure 11: Example images and annotations from the BKAI-IGH NeoPolyp-Small dataset

The annotations in Figure 11 serve as ground truth for Dice Score evaluation and for analyzing whether post-processing improves mask consistency. The three-class setting is important because a mask is not only a geometric localization result; it also carries semantic information

that may be used by the VQA agent. Therefore, segmentation errors can affect both direct mask interpretation and downstream visual grounding.

7.4 Kvasir-VQA-X1 Dataset

Kvasir-VQA-X1 is used as the primary dataset for LoRA fine-tuning and evaluation of the Polyp VQA VLM. It is a gastrointestinal endoscopy visual question answering dataset organized as image-question-answer triples [36]. The public dataset extends Kvasir-VQA with 159,549 new question-answer pairs designed for medical reasoning and robustness evaluation. In the local training split used by this thesis, the full VQA setting contains 6,212 original endoscopic images and 143,594 image-question-answer samples.




Figure 12: Example structure of the Kvasir-VQA-X1 dataset for gastrointestinal endoscopy VQA

This structure is suitable for the thesis because the model must connect endoscopic visual evidence with a natural-language question rather than only predict an image label or segmentation mask. Kvasir-VQA-X1 is used in two experimental formats:

- **Original-image VQA:** the model receives the endoscopic image and the question;
- **Visual-grounded VQA:** the model receives the original image, a segmentation mask or overlay, and the question.

The first format evaluates direct visual reasoning from the original frame. The second evaluates whether segmentation-derived grounding improves downstream VQA by giving the model an explicit cue about the region of interest.

The full Kvasir-VQA-X1 training split contains 6,212 original endoscopic images and 143,594 image-question-answer samples. However, segmentation-based visual grounding can only be applied when an image has a matched segmentation mask. After filtering for mask availability, 1,336 images remain with usable segmentation masks, corresponding to 42,382 training samples. Table 6 summarizes this filtering step.

Training Data Setting	Number of Images	Number of Samples
Original Kvasir-VQA-X1 training data	6,212	143,594
Mask-filtered visual grounding subset	1,336	42,382

Table 6: Kvasir-VQA-X1 training data before and after segmentation-mask filtering

This filtering step is necessary because visual grounding requires both the original image and a corresponding segmentation mask. Therefore, improvements from visual grounding should be compared against a LoRA-only baseline evaluated on the same mask-matched subset, not against a model evaluated on the full dataset. Evaluation data are kept separate from training data to avoid image-level or question-level leakage, and preprocessing is kept conservative to preserve clinically relevant visual evidence.

7.5 Polyp VQA Fine-tuning Setup

The Polyp VQA models are fine-tuned with LoRA. The training objective is to generate correct and clinically relevant answers from image-question inputs. The fine-tuning setup compares two base-model families, MedGemma-4B-it and Qwen3.5-4B, under both original-image input and visual-grounded input.

Table 7 summarizes the main hyperparameters used for LoRA fine-tuning. The same setup is used for MedGemma-4B-it and Qwen3.5-4B in the original-image setting, with the corresponding base model changed according to the experiment.

Hyperparameter	Original-image LoRA	Visual-grounded LoRA
Training dataset	Kvasir-VQA-x1 train	Mask-filtered Kvasir-VQA-x1 train
Validation dataset	Kvasir-VQA-x1 test-1000	Mask-matched validation subset
Base models	MedGemma-4B-it, Qwen3.5-4B	MedGemma-4B-it
Maximum sequence length	512	1024
Torch dtype	bfloat16	bfloat16
Number of epochs	1	3
Train batch size per device	4	4
Evaluation batch size per device	4	4
Gradient accumulation steps	4	4
Learning rate	2×10^{-5}	2×10^{-5}
Learning-rate scheduler	linear	linear
Warmup ratio	0.03	0.03
Weight decay	0.01	0.01
LoRA rank	16	16
LoRA alpha	64	64
Freeze vision encoder	true	true
Gradient checkpointing	true	true
Best-model metric	eval token accuracy	eval token accuracy
Checkpoint save interval	1000 steps	1000 steps
Maximum saved checkpoints	2	2
Logging interval	20 steps	20 steps
Data loading workers	2	2
Dataset preprocessing processes	2	2
Experiment tracking	Weights & Biases	Weights & Biases

Table 7: LoRA fine-tuning hyperparameters for Polyp VQA

The visual-grounded setting uses a longer maximum sequence length of 1024 because each example contains two visual inputs: the original endoscopic image and the segmentation-derived grounding image. In contrast, the original-image LoRA setting uses a maximum length of 512 because it only includes one image and the question-answer text.

The number of epochs is also different between the two settings. Original-image LoRA is trained for one epoch on the full training set, while the visual-grounded setting is trained for three epochs on the filtered subset. Since the mask-filtered subset is substantially smaller than the full VQA training set, the additional epochs are used to keep the number of training iterations more comparable between the two training strategies. This reduces optimization-step

count as a potential confounding variable when comparing original-image fine-tuning and visual-grounded fine-tuning. Training and validation loss curves did not show signs of overfitting across the visual-grounded epochs, supporting the use of additional epochs on the smaller subset.

The original-image setting trains each model to answer from the endoscopic frame and question only. The visual-grounded setting extends the input with segmentation-derived evidence, such as a mask or overlay. To make the comparison fair, each base model should be evaluated under both input settings with the same LoRA strategy and the same evaluation split. This separates two effects: the difference between base-model backbones and the effect of explicit segmentation guidance.

The output format should be consistent across training and evaluation. Short factual answers are easier to evaluate with lexical metrics, while longer explanatory answers may better reflect clinical reasoning but are harder to score with BLEU or ROUGE. For this reason, LLM Judge is included as a complementary evaluation method.

7.6 Memory Evaluation Scenarios

The memory layer should be evaluated using multi-session Vietnamese medical scenarios. Representative scenarios include:

- a user reporting a medication allergy in one session and asking about a drug later;
- a user with a chronic condition asking follow-up symptom questions;
- repeated symptoms across sessions;
- medication use that should influence later advice;
- follow-up questions after medical image analysis.

7.7 Evaluation Metrics

Different components require different metrics.

Text Module. One-Hop Accuracy and Multi-Hop Awareness Score measure direct reasoning and awareness of missing information.

Segmentation. Dice Score measures overlap between predicted masks and ground-truth masks.

Polyp VQA. BLEU, ROUGE-L, and LLM Judge evaluate answer quality. BLEU and ROUGE-L measure textual overlap, while LLM Judge evaluates semantic correctness, completeness, visual grounding, and safety.

Memory Layer. Memory retrieval relevance, response consistency, repeated-question reduction, routing accuracy, and latency can be used to evaluate the impact of long-term memory.

7.8 Experimental Protocol

The experimental protocol follows a component-wise and system-wise evaluation design. Component-wise evaluation measures each technical module with metrics suitable for its output format. The Text Module is evaluated with one-hop and multi-hop reasoning benchmarks. The Segmentation Agent is evaluated with Dice Score. The Polyp VQA Agent is evaluated on Kvasir-VQA-X1 using lexical and judge-based metrics. The Memory Layer is evaluated through multi-session scenarios.

System-wise evaluation focuses on interaction behavior. This includes whether the Decision Agent selects the correct processing path, whether the Text Module asks for missing information, whether visual grounding improves or harms VQA, and whether memory improves continuity without injecting irrelevant context. This distinction is important because a multi-agent assistant can fail even when individual components perform reasonably well. For example,

a good segmentation model does not help if the image query is routed to the wrong agent.

For all experiments, qualitative analysis should accompany numerical metrics. Medical AI outputs are difficult to evaluate with a single number because correctness, safety, clarity, and uncertainty handling are all important. The report should therefore include representative successes and failures, especially for high-risk symptom questions, incorrect retrieval cases, segmentation failures, and VQA examples where the mask changes the answer.

8 Chapter 8. Results and Analysis

8.1 Text Module Results and Analysis

The Text Module is evaluated using one-hop and multi-hop medical reasoning benchmarks. Table 8 summarizes the comparison.

Model Variant	One-Hop Accuracy	Multi-Hop Awareness Score
KG + MedLLM	86.95%	70.83%
RAG + MedLLM	67.97%	40.38%
MedLLM	60.87%	50.72%

Table 8: Text Module benchmark results

The KG + MedLLM setting achieves the best performance in both One-Hop Accuracy and Multi-Hop Awareness Score. This suggests that structured medical relations help the model reason more accurately when the task depends on disease-symptom links or explicit medical relationships.

RAG + MedLLM performs better than the base model on one-hop reasoning but is weaker on Multi-Hop Awareness. This indicates that retrieved text is useful for explanation, but it may not always help the model recognize when information is incomplete. The base MedLLM setting lacks external grounding and therefore performs worse on structured reasoning.

The difference between One-Hop Accuracy and Multi-Hop Awareness is important. One-Hop Accuracy measures whether the system can answer when the required relation is available. Multi-Hop Awareness measures whether the system understands that a direct answer may be unsafe when key information is missing. A medical assistant should perform well on both dimensions. High answer accuracy alone is not sufficient if the model answers prematurely in incomplete clinical scenarios.

The stronger performance of KG + MedLLM suggests that explicit relations help the model identify clinically relevant links more reliably than text retrieval alone. However, the result should not be interpreted as showing that KG is always superior to RAG. KG is stronger for structured relation reasoning, while RAG remains useful for explanatory responses and broader medical education. The practical conclusion is that the two sources should be coordinated rather than treated as mutually exclusive.

The lower Multi-Hop Awareness Score of RAG + MedLLM suggests that adding retrieved documents may encourage answer generation even when the user query lacks enough clinical detail. This is a common risk in retrieval-augmented systems: retrieved evidence can make the model appear grounded while the original question remains underspecified. The information sufficiency step is therefore necessary before retrieval-based answer synthesis.

The base MedLLM result shows that a language model alone can answer some medical questions but performs worse when structured external knowledge is required. This supports the design choice of placing the LLM inside a retrieval and graph-augmented pipeline rather than using it as the only reasoning source.

8.2 Polyp Segmentation Results and Analysis

Table 9 presents Dice Score results for several segmentation models with and without post-processing.

Models	DICE Score	
	Without Processing	With Processing
PraNet	80.02%	82.94%
DeepLabV3+ (RegNetx320)	82.56%	86.33%
DeepLabV3+ (ResNet50)	81.12%	85.62%
EMCAD (PVT2-b2)	80.93%	83.76%
U-Net	75.95%	-
U-Net++	77.54%	-

Table 9: Dice Score comparison for polyp segmentation models

Post-processing improves segmentation consistency across the evaluated models. The improvement is caused by reducing fragmented and mixed-color predictions inside connected regions. DeepLabV3+ with RegNetx320 achieves the strongest result after post-processing, while DeepLabV3+ with ResNet50 also shows a large improvement.

This confirms that post-processing is useful even when the base segmentation model is strong. In a medical assistant, a cleaner mask is valuable not only for visualization but also for visual grounding in Polyp VQA.

The first stage, connected component analysis, changes the interpretation of the mask from isolated pixel predictions to region-level candidates. This is important because a polyp is normally interpreted as a spatial object rather than as independent pixels. By grouping all adjacent non-background pixels into connected regions, the method separates foreground structures and makes it possible to refine each candidate lesion independently. Small noisy islands can also be inspected or filtered in later extensions, although the current thesis focuses on region-wise label refinement.

The second stage, region-wise color dominance refinement, addresses a common error in multi-class segmentation: a single connected lesion may contain scattered red and green pixels even when one class clearly dominates. In the BKAI-IGH NeoPolyp setting, this is problematic because the model must distinguish neoplastic and non-neoplastic regions. If a connected region is visually interpreted as one lesion but the predicted class labels are internally fragmented, the mask becomes harder to use for both human review and VQA grounding. Assigning the dominant class to the whole connected component reduces this inconsistency.

This post-processing strategy is different from simply smoothing the mask. It preserves the spatial extent of each detected connected region, but regularizes the semantic label inside the region. The dominance ratio ρ_i explains when this operation is reliable. If ρ_i is close to 1, the model has already made a strong region-level class decision and the refinement removes only a small amount of minority-label noise. If ρ_i is close to 0.5, the region contains competing class evidence, and forcing one label may be less reliable. Therefore, the improvement comes mainly from correcting high-dominance intra-region label noise rather than from changing the model’s localization ability.

The qualitative examples in Figure 8 align with the quantitative trend in Table 9. The largest gains appear in settings where the raw model prediction is close to the correct region but contains class fragments. In those cases, $C_k(R_i)$ identifies the class that dominates the connected component, and the update rule for $M'(x, y)$ increases overlap with the ground-truth class by removing minority-class artifacts. This effect is especially visible in multi-class segmentation because a pixel can be spatially inside the lesion but still counted as wrong if its

neoplastic/non-neoplastic label is incorrect.

Post-processing should therefore be evaluated as part of the complete segmentation pipeline. A raw neural network output may have good average overlap but still produce fragmented regions that are difficult for users or downstream agents to interpret. Connected component analysis and color dominance refinement improve spatial coherence, which is especially important when the mask is later passed to a VQA model as visual evidence.

The improvement after post-processing is also relevant to human review. A cleaner overlay is easier to inspect, which supports the human-in-the-loop design of the system. If the mask is visually noisy, reviewers may spend more effort determining whether the highlighted region is meaningful. Therefore, post-processing contributes to both quantitative overlap and practical interpretability.

The method also has clear limitations. If two touching regions actually belong to different semantic classes, dominance refinement may incorrectly force them into one class. If the segmentation model misses a lesion, the post-processing step cannot recover it because it only operates on predicted foreground regions where $I(x, y) = 1$. If the model detects background artifacts as foreground, the connected component step may preserve those false regions unless an additional size, shape, or confidence filter is introduced. For this reason, post-processing should be interpreted as a consistency refinement module, not as a replacement for the segmentation model.

For the downstream Polyp VQA setting, the main benefit is that the overlay becomes a clearer grounding signal. A fragmented overlay can mislead the VLM by presenting contradictory visual cues inside the same lesion. A refined overlay is easier to align with questions about location, suspicious regions, and neoplastic versus non-neoplastic interpretation. However, the same dependency also creates a risk: if dominance refinement assigns the wrong class to a region, the VQA model may produce a more confident but incorrect answer. This motivates the thesis design of evaluating original-image VQA and segmentation-grounded VQA separately.

8.3 Polyp VQA Fine-tuning Results and Analysis on Kvasir-VQA-X1

The Polyp VQA VLM is evaluated using BLEU-4, ROUGE-L, and an LLM-as-a-judge protocol. BLEU-4 and ROUGE-L measure lexical overlap with the reference answer, while the judge-based score is used to assess semantic and clinical answer quality. The judge model is `google/medgemma-27b-it`. For each prediction, the judge receives the question, ground-truth answer, and model prediction, and returns four scores from 1 to 5: factual accuracy, completeness, clinical relevance, and no hallucination. The final weighted score is computed as:

$$\begin{aligned} \text{Score} = & 0.40 \times \text{Factual Accuracy} + 0.40 \times \text{Completeness} \\ & + 0.10 \times \text{Clinical Relevance} + 0.10 \times \text{No Hallucination}. \end{aligned}$$

This weighting gives the highest importance to factual correctness and coverage of the reference answer, while still accounting for clinical language quality and hallucination control. This is appropriate for medical VQA because an answer that is fluent but factually wrong or incomplete should not be considered reliable.

Model Setting	Samples	BLEU-4	ROUGE-L	LLM Judge
MedGemma-4B-it base	1000	0.0055	0.1154	2.7521
MedGemma-4B-it + LoRA fine-tuning	1000	0.3657	0.6936	4.6351
Qwen3.5-4B base	1000	0.0261	0.2078	3.4720
Qwen3.5-4B + LoRA fine-tuning	1000	0.1935	0.5211	4.5341
MedGemma-4B-it + LoRA only, mask-matched subset	305	–	–	4.4269
MedGemma-4B-it + LoRA + visual grounding, mask-matched subset	305	0.4139	0.6966	4.5751

Table 10: Polyp VQA evaluation on Kvasir-VQA-X1

The results show that LoRA fine-tuning is the main driver of performance improvement. MedGemma-4B-it improves from a BLEU-4 score of 0.0055 and LLM Judge score of 2.7521 to 0.3657 and 4.6351, respectively. Qwen3.5-4B also benefits strongly from LoRA, increasing from 0.0261 BLEU-4 and 3.4720 judge score to 0.1935 and 4.5341. These improvements indicate that the base VLMs do not directly match the answer style and clinical content distribution of Kvasir-VQA-X1, while lightweight task adaptation substantially improves both lexical matching and semantic answer quality.

Between the two backbones, MedGemma-4B-it achieves stronger fine-tuned performance than Qwen3.5-4B. This is consistent with the expectation that a medically oriented VLM backbone is better aligned with endoscopic visual concepts and clinical answer generation. Qwen3.5-4B remains a strong baseline after fine-tuning, but its lower BLEU-4 and ROUGE-L suggest that it may be less aligned with the dataset’s medical visual-language distribution.

The visual grounding comparison must be interpreted on the mask-matched subset only. As shown in Table 6, only 1,336 of the 6,212 training images have usable segmentation masks, reducing the training samples from 143,594 to 42,382 for the grounded setting. Therefore, the fair comparison is not between the 1000-sample MedGemma LoRA result and the visual-grounded result. Instead, the correct baseline is the MedGemma LoRA-only model evaluated again on the same 305 evaluation samples that have matched visual grounding inputs. Under this controlled comparison, visual grounding improves the LLM Judge score from 4.4269 to 4.5751. This suggests that segmentation masks provide useful spatial evidence for VQA when the mask is available and aligned with the question.

The original-image setting evaluates whether the VLM can identify relevant visual evidence without additional guidance. The visual-grounded setting evaluates whether a segmentation mask improves answer quality by focusing attention on the polyp region. Since only a subset of samples contains matched segmentation masks, the visual grounding analysis is conducted on 305 mask-matched samples.

LLM Judge Criterion	MedGemma LoRA only	MedGemma LoRA + visual grounding	Difference
Factual Accuracy	4.2918	4.4951	+0.2033
Completeness	4.3607	4.5016	+0.1409
Clinical Relevance	4.8984	4.9246	+0.0262
No Hallucination	4.7607	4.8393	+0.0786
Weighted Score	4.4269	4.5751	+0.1482
Samples in 4.0–5.0 range	238/305	255/305	+17

Table 11: LLM Judge comparison between LoRA-only and visual-grounded MedGemma on the mask-matched subset

The criterion-level judge results in Table 11 indicate that visual grounding improves all four judge criteria. The largest gain appears in factual accuracy (+0.2033), followed by completeness (+0.1409). This pattern is important: the mask does not merely make the generated answer

more fluent, but appears to help the model select more correct visual evidence and include more of the expected clinical content. The smaller gain in clinical relevance (+0.0262) is also reasonable because the LoRA-only model already has a very high clinical relevance score of 4.8984, leaving little room for improvement.

The hallucination score also improves from 4.7607 to 4.8393. This suggests that explicit spatial grounding may reduce the tendency to add unsupported visual claims. In medical VQA, this is valuable because a confident but visually unsupported answer can be unsafe even when it uses correct clinical terminology. The distribution of weighted scores also improves: the number of answers in the 4.0–5.0 range increases from 238 to 255 out of 305 samples. Therefore, the improvement is not only reflected in the average score, but also in a larger number of high-quality answers.

These results support the design choice of connecting the Segmentation Agent with the Polyp VQA Agent. The segmentation mask acts as an intermediate visual evidence signal, allowing the VQA model to focus on the lesion area instead of relying only on global image context or dataset priors. However, the result should not be interpreted as proof that visual grounding is always beneficial. The grounded setting depends on mask availability and mask quality. If the segmentation model misses the lesion, highlights irrelevant tissue, or assigns an incorrect region, the same grounding mechanism may guide the VQA model toward the wrong evidence. For this reason, visual grounding should be treated as a useful but conditional intervention, and future evaluation should include failure cases where incorrect masks degrade the answer.

9 Chapter 9. Overall Discussion and Limitations

9.1 Strengths of the Proposed System

The main strength of the system is its modular AI architecture. Instead of relying on a single model, the system combines specialized agents for text reasoning, graph reasoning, retrieval, visual analysis, VQA, memory, and validation. This design improves extensibility and makes the reasoning process easier to analyze.

The architecture also supports complementary reasoning sources. RAG provides explanatory text, KG provides structured relations, Web Search provides fallback coverage, and memory provides patient-specific continuity. In the Vision Module, segmentation provides localized evidence while the VQA agent converts visual and textual input into an answer. This combination allows the system to handle a wider range of medical-assistant tasks than a single text-only chatbot.

Another strength is that the system makes several uncertainty points explicit. Information sufficiency checks identify incomplete symptom descriptions. Routing determines whether a specialized agent is appropriate. Visual grounding is evaluated as a separate input setting rather than assumed to be beneficial. Memory is treated as patient-reported context rather than verified diagnosis. These design choices improve transparency and make the system easier to evaluate.

9.2 Text Module Limitations

The Text Module depends on the coverage and quality of its data sources. RAG may retrieve irrelevant documents, and the Knowledge Graph may not contain all relevant medical relations. Multi-hop QA also depends on correctly identifying missing information.

Vietnamese medical language adds additional challenges. Users may describe symptoms in informal language, use regional terms, or mix Vietnamese with English medical terminology. If query expansion does not capture these variations, retrieval quality may decrease. If expansion

is too broad, the system may retrieve clinically unrelated contexts. This trade-off affects both answer accuracy and safety.

The Text Module also cannot guarantee clinical correctness. Even when KG and RAG provide useful evidence, the final synthesis is still generated by an LLM and may omit caveats or overstate uncertainty. Human validation and expert review remain necessary for high-risk cases.

9.3 Vision Module Limitations

The Vision Module is limited by the quality of segmentation and VQA fine-tuning data. General medical image interpretation is broad and difficult to validate across all modalities. Polyp VQA is more focused but depends on the relevance of the training dataset and the quality of segmentation masks.

Segmentation errors can propagate to downstream VQA. If the mask misses the polyp region, highlights irrelevant tissue, or assigns the wrong class, the grounded VQA input may mislead the model. Therefore, visual grounding is not automatically safer than original-image VQA. It should be evaluated as an intervention that can help in some cases and harm in others.

The General Medical Image Agent has broader coverage but weaker guarantees. Without modality-specific fine-tuning and evaluation, its outputs should remain descriptive and cautious. This agent is useful for system extensibility, but the thesis does not claim comprehensive diagnostic performance across all medical image types.

9.4 Kvasir-VQA-X1 Limitations and Domain Transfer

Kvasir-VQA-X1 is useful for endoscopic VQA, but dataset coverage and annotation style may not fully match real deployment images or user questions. Metrics such as BLEU and ROUGE-L may not capture clinical correctness, so LLM Judge and qualitative analysis are also needed.

Domain transfer is a central concern. A model fine-tuned on Kvasir-VQA-X1 may learn the style and distribution of that dataset but still struggle with images from different devices, lighting conditions, bowel preparation quality, or clinical populations. User questions in deployment may also be longer, less structured, and more ambiguous than dataset questions. These differences should be considered when interpreting evaluation results.

The two VQA settings also introduce different risks. The original-image setting depends on the VLM’s own attention and visual reasoning. The grounded setting depends on both segmentation quality and the VLM’s ability to use the overlay correctly. A fair discussion must therefore analyze not only which setting scores higher but also why it succeeds or fails.

9.5 Memory Layer Limitations

Memory can be stale, incomplete, or incorrectly extracted. Patient-reported facts may change over time. Therefore, memory should be used as supporting context only and should not override current user input or medical evidence.

A memory layer can also introduce privacy and governance concerns. Even clinically useful facts may be sensitive. The system should avoid storing unnecessary personally identifiable information and should support correction or deletion of stored facts. In the current thesis scope, memory is evaluated as a system integration feature rather than as a complete clinical record management system.

Another limitation is that memory retrieval may surface irrelevant old facts. If the response generator uses such facts without checking relevance, personalization can become distracting or

unsafe. The agent should only use memory when it is directly related to the current question.

9.6 Evaluation and Safety Limitations

The evaluation is component-oriented and system-oriented, but it is not a clinical trial. The system should be positioned as medical decision support, not as a replacement for physicians. More rigorous expert evaluation is necessary before any clinical use.

The metrics used in this thesis capture important but incomplete aspects of performance. Dice Score measures segmentation overlap but does not fully measure clinical usefulness. BLEU and ROUGE-L measure text similarity but not necessarily medical correctness. LLM Judge can evaluate semantic quality but may contain evaluator bias. One-Hop Accuracy and Multi-Hop Awareness Score measure reasoning behavior but cannot cover all clinical scenarios.

For these reasons, the results should be interpreted as evidence of technical feasibility and comparative system behavior, not as proof of clinical readiness. Future evaluation should include medical expert review, larger datasets, prospective testing, and safety analysis under realistic user interactions.

10 Chapter 10. Conclusion and Future Work

10.1 Summary of Achievements

This thesis presents a multi-agent medical assistant that combines text reasoning, image analysis, long-term memory, and human-in-the-loop validation. The Text Module integrates Multi-hop QA, RAG, Knowledge Graph reasoning, and fallback search. The Vision Module integrates General Medical Image Agent, Polyp Segmentation Agent, and Polyp VQA Agent. The Memory Layer provides long-term patient context through Mem0.

10.2 Technical Contributions

The technical contributions include:

- a graph-orchestrated multi-agent architecture for medical reasoning;
- a Text Module combining structured and unstructured medical knowledge;
- a polyp segmentation pipeline with post-processing and overlay generation;
- LoRA fine-tuned Polyp VQA models based on MedGemma-4B-it and Qwen3.5-4B, evaluated with original-image and visual-grounded settings;
- a long-term memory layer for patient-context continuity;
- a human-in-the-loop validation mechanism for feedback and review.

10.3 Future Work

Future work should focus on expanding Vietnamese medical datasets, improving Knowledge Graph coverage, strengthening RAG retrieval and reranking, improving visual grounding robustness, refining LoRA fine-tuning for VQA, adding expert validation, and developing stronger memory governance mechanisms. Clinical evaluation with medical professionals is also necessary before real-world deployment.

References

- [1] LangGraph Documentation - Build language agents as graphs. <https://langchain-ai.github.io/langgraph/>
- [2] Google AI. Gemini 2.0 Flash: A new generation of multimodal AI. <https://deepmind.google/technologies/gemini/flash/>
- [3] Qdrant Documentation - Vector Database for AI Applications. <https://qdrant.tech/documentation/>
- [4] Tavily AI - Real-time search API for AI agents. <https://tavily.com/>
- [5] Wang, L., Yang, N., Huang, X., et al. Text Embeddings by Weakly-Supervised Contrastive Pre-training. *arXiv preprint arXiv:2212.03533*, 2022.
- [6] Nguyen, T.M., Le, V.H., Phan, X.H. VimedAQA: A Vietnamese Medical Question Answering Dataset. *Proceedings of the 12th International Conference on Knowledge and Systems Engineering (KSE)*, 2023.
- [7] Lewis, P., Perez, E., Piktus, A., et al. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. *Advances in Neural Information Processing Systems*, 2020.
- [8] Yu, K.H., Beam, A.L., Kohane, I.S. Artificial intelligence in healthcare. *Nature Biomedical Engineering*, vol. 2, no. 10, pp. 719-731, 2018.
- [9] Zhang, Y., Chen, H., Song, Y., et al. Multimodal machine learning for medical image analysis: A comprehensive review. *Medical Image Analysis*, vol. 78, 102384, 2022.
- [10] Singhal, K., Azizi, S., Tu, T., et al. Large language models encode clinical knowledge. *Nature*, vol. 620, pp. 172-180, 2023.
- [11] Singhal, K., Tu, T., Gottweis, J., et al. Towards Expert-Level Medical Question Answering with Large Language Models. *arXiv preprint arXiv:2305.09617*, 2023.
- [12] Wu, C., Lin, W., Zhang, X., Zhang, Y., Wang, Y., Xie, W. PMC-LLaMA: Towards Building Open-source Language Models for Medicine. *arXiv preprint arXiv:2304.14454*, 2023.
- [13] Labrak, Y., Bazoge, A., Morin, E., Gourraud, P.A., Rouvier, M., Dufour, R. BioMistral: A Collection of Open-Source Pretrained Large Language Models for Medical Domains. *arXiv preprint arXiv:2402.10373*, 2024.
- [14] Li, C., Wong, C., Zhang, S., et al. LLaVA-Med: Training a Large Language-and-Vision Assistant for Biomedicine in One Day. *Advances in Neural Information Processing Systems*, 2023.
- [15] Moor, M., Huang, Q., Wu, S., et al. Med-Flamingo: a Multimodal Medical Few-shot Learner. *arXiv preprint arXiv:2307.15189*, 2023.
- [16] Tonmoy, S.M.T.I., Zaman, S.M.M., Jain, V., et al. A Comprehensive Survey of Hallucination Mitigation Techniques in Large Language Models. *arXiv preprint arXiv:2401.01313*, 2024.
- [17] Jin, Q., Dhingra, B., Liu, Z., Cohen, W.W., Lu, X. PubMedQA: A Dataset for Biomedical Research Question Answering. *Proceedings of EMNLP-IJCNLP*, 2019.

- [18] Jin, D., Pan, E., Oufattole, N., Weng, W.H., Fang, H., Szolovits, P. What Disease does this Patient Have? A Large-scale Open Domain Question Answering Dataset from Medical Exams. *Applied Sciences*, vol. 11, no. 14, 2021.
- [19] Pal, A., Umapathi, L.K., Sankarasubbu, M. MedMCQA: A Large-scale Multi-Subject Multi-Choice Dataset for Medical Domain Question Answering. *Conference on Health, Inference, and Learning*, 2022.
- [20] Yang, Z., Qi, P., Zhang, S., et al. HotpotQA: A Dataset for Diverse, Explainable Multi-hop Question Answering. *Proceedings of EMNLP*, 2018.
- [21] Liu, X., Faes, L., Kale, A.U., et al. A comparison of deep learning performance against health-care professionals in detecting diseases from medical imaging. *Nature Medicine*, vol. 25, no. 8, pp. 1252-1258, 2019.
- [22] Bickmore, T.W., Gruber, A., Picard, R. Establishing the computer-patient working alliance in automated health behavior change interventions. *Patient Education and Counseling*, vol. 59, no. 1, pp. 21-30, 2005.
- [23] Nguyen, D.Q., Nguyen, A.T. PhoBERT: Pre-trained language models for Vietnamese. *Findings of EMNLP*, 2020.
- [24] Laranjo, L., Dunn, A.G., Tong, H.L., et al. Conversational agents in healthcare: a systematic review. *Journal of the American Medical Informatics Association*, vol. 25, no. 9, pp. 1248-1258, 2018.
- [25] Tampuu, A., Matiisen, T., Kodelja, D., et al. Multiagent cooperation and competition with deep reinforcement learning. *PloS one*, vol. 12, no. 4, e0172395, 2017.
- [26] Nogueira, R., Cho, K. Passage Re-ranking with BERT. *arXiv preprint arXiv:1901.04085*, 2019.
- [27] Ronneberger, O., Fischer, P., Brox, T. U-Net: Convolutional Networks for Biomedical Image Segmentation. *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, 2015.
- [28] Zhou, Z., Siddiquee, M.M.R., Tajbakhsh, N., Liang, J. UNet++: A Nested U-Net Architecture for Medical Image Segmentation. *Deep Learning in Medical Image Analysis and Multimodal Learning for Clinical Decision Support*, pp. 3-11, 2018.
- [29] Chen, L.C., Zhu, Y., Papandreou, G., Schroff, F., Adam, H. Encoder-Decoder with Atrous Separable Convolution for Semantic Image Segmentation. *European Conference on Computer Vision (ECCV)*, 2018.
- [30] Fan, D.P., Ji, G.P., Zhou, T., Chen, G., Fu, H., Shen, J., Shao, L. PraNet: Parallel Reverse Attention Network for Polyp Segmentation. *International Conference on Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, 2020.
- [31] Jha, D., Smedsrud, P.H., Riegler, M.A., Halvorsen, P., de Lange, T., Johansen, D., Johansen, H.D. Kvasir-SEG: A Segmented Polyp Dataset. *International Conference on Multimedia Modeling*, 2020.
- [32] Dong, B., Wang, W., Fan, D.P., Li, J., Fu, H., Shao, L. Polyp-PVT: Polyp Segmentation with Pyramid Vision Transformers. *arXiv preprint arXiv:2108.06932*, 2021.

- [33] Rahman, M.M., Munir, M., Marculescu, R. EMCAD: Efficient Multi-scale Convolutional Attention Decoding for Medical Image Segmentation. *arXiv preprint arXiv:2405.06880*, 2024.
- [34] Li, Y., Hu, M., Yang, X. Polyp-SAM: Transfer SAM for Polyp Segmentation. *arXiv preprint arXiv:2305.00293*, 2023.
- [35] Gautam, S., Storås, A., Midoglu, C., Hicks, S.A., Thambawita, V., Halvorsen, P., Riegler, M.A. Kvasir-VQA: A Text-Image Pair GI Tract Dataset. *arXiv preprint arXiv:2409.01437*, 2024.
- [36] Gautam, S., Riegler, M.A., Halvorsen, P. Kvasir-VQA-x1: A Multimodal Dataset for Medical Reasoning and Robust MedVQA in Gastrointestinal Endoscopy. *arXiv preprint arXiv:2506.09958*, 2025.
- [37] Hu, E.J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., Chen, W. LoRA: Low-Rank Adaptation of Large Language Models. *International Conference on Learning Representations (ICLR)*, 2022.
- [38] Sellergren, A., Kazemzadeh, S., Jaroensri, T., et al. MedGemma Technical Report. *arXiv preprint arXiv:2507.05201*, 2025.
- [39] Yang, A., Li, A., Yang, B., et al. Qwen3 Technical Report. *arXiv preprint arXiv:2505.09388*, 2025.